

Under the Covers

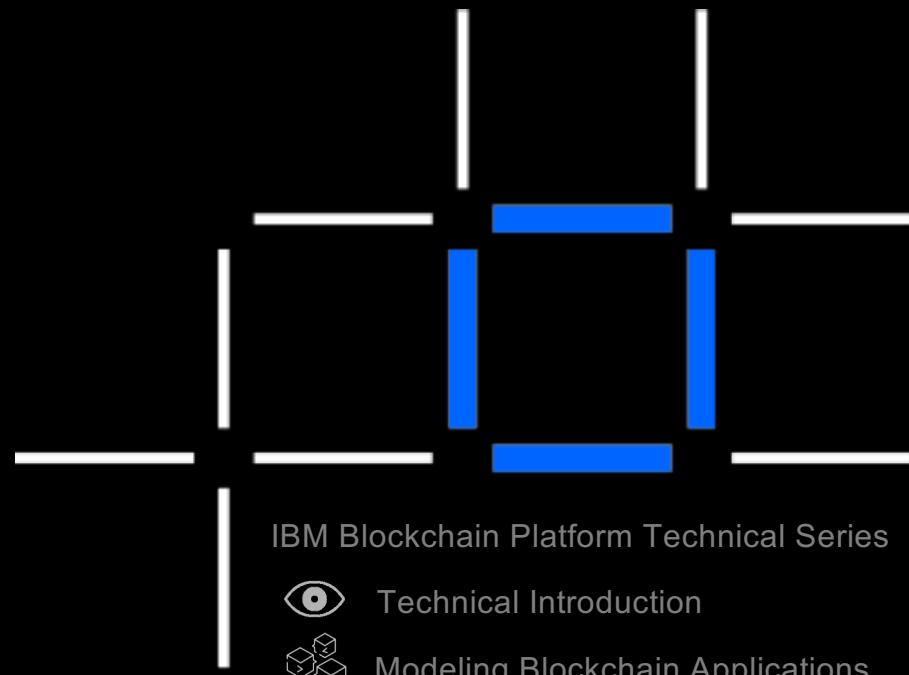
A Technical Deep-Dive on Hyperledger Fabric

Olivier VALLOD

IT Architect

IBM Client Center Montpellier, France

olivier.vallod@fr.ibm.com



IBM Blockchain Platform Technical Series

👁️ Technical Introduction

📦 Modeling Blockchain Applications

🧩 **Under the Covers**

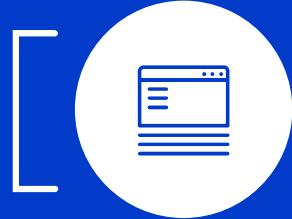
🏗️ Architectural Good Practices

⊕ What's New in Technology

V4.08, 24 September 2019

IBM Blockchain





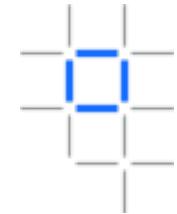
Project Status and
Roadmap



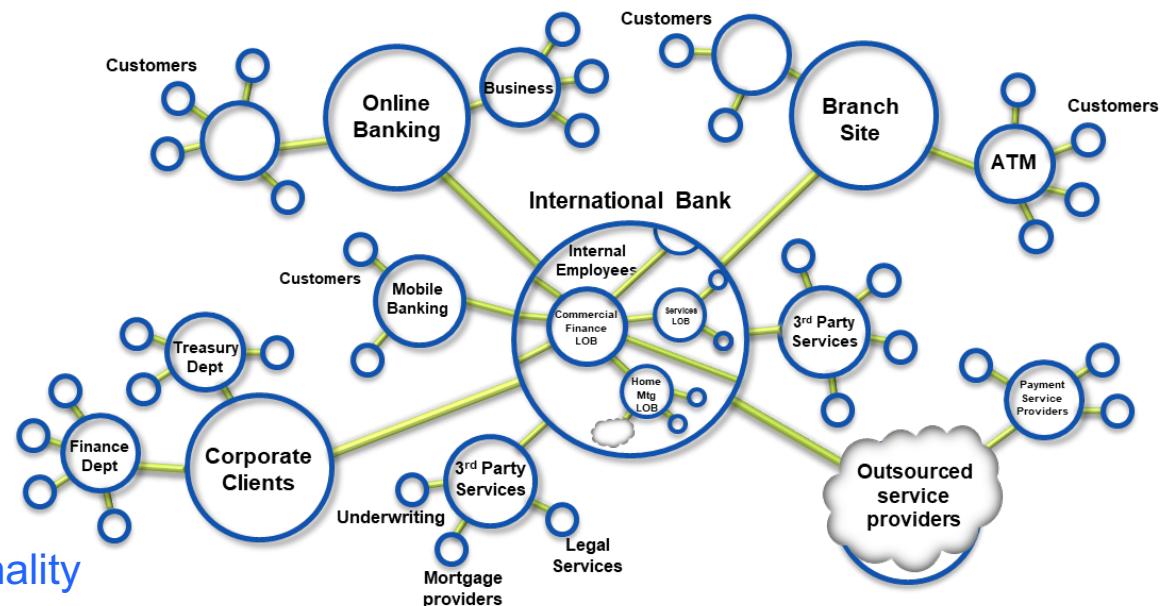
Hyperledger Fabric
Technical Dive



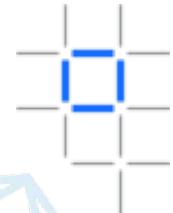
Blockchain Recap



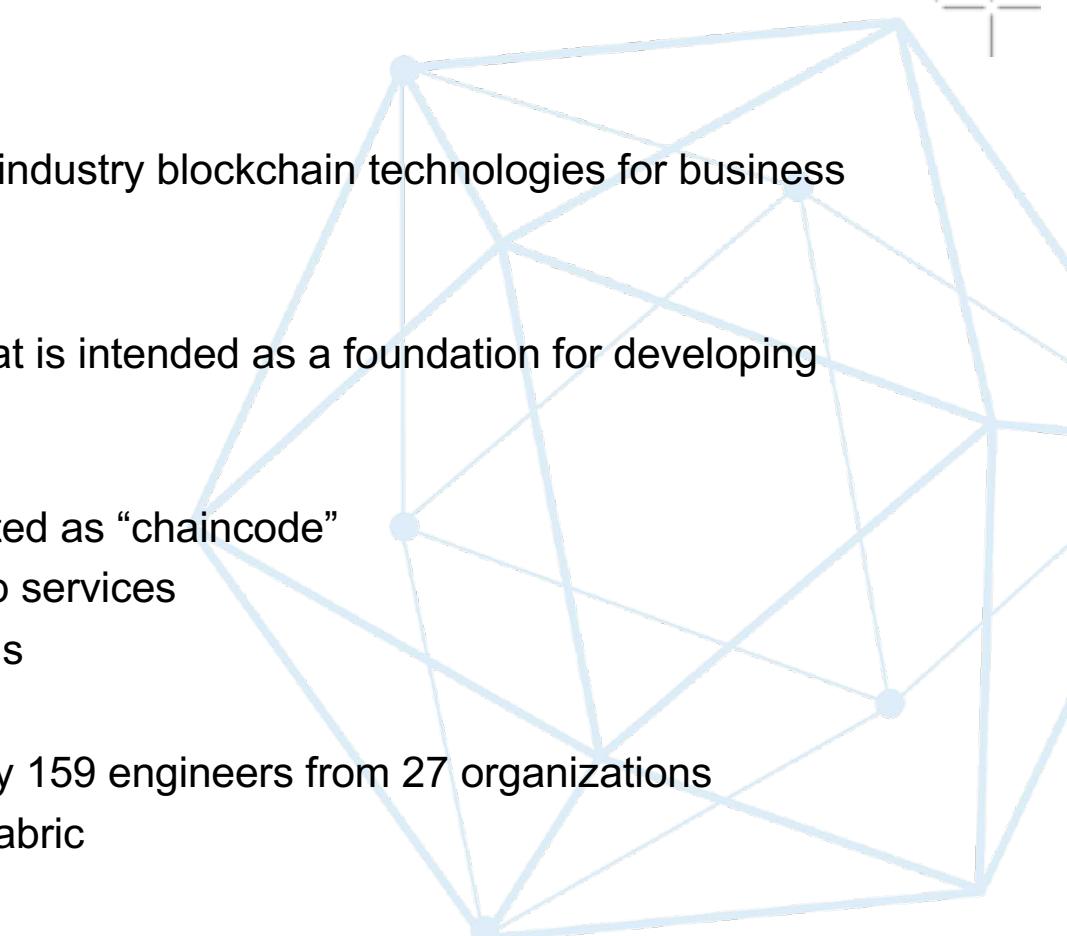
- Blockchain builds on basic business concepts
 - **Business Networks** connect businesses
 - **Participants** with **Identity**
 - **Assets** flow over business networks
 - **Transactions** describe asset exchange
 - **Contracts** underpin transactions
 - The **ledger** is a log of transactions
- Blockchain is a shared, replicated ledger
 - **Consensus**, **provenance**, **immutability**, **finality**



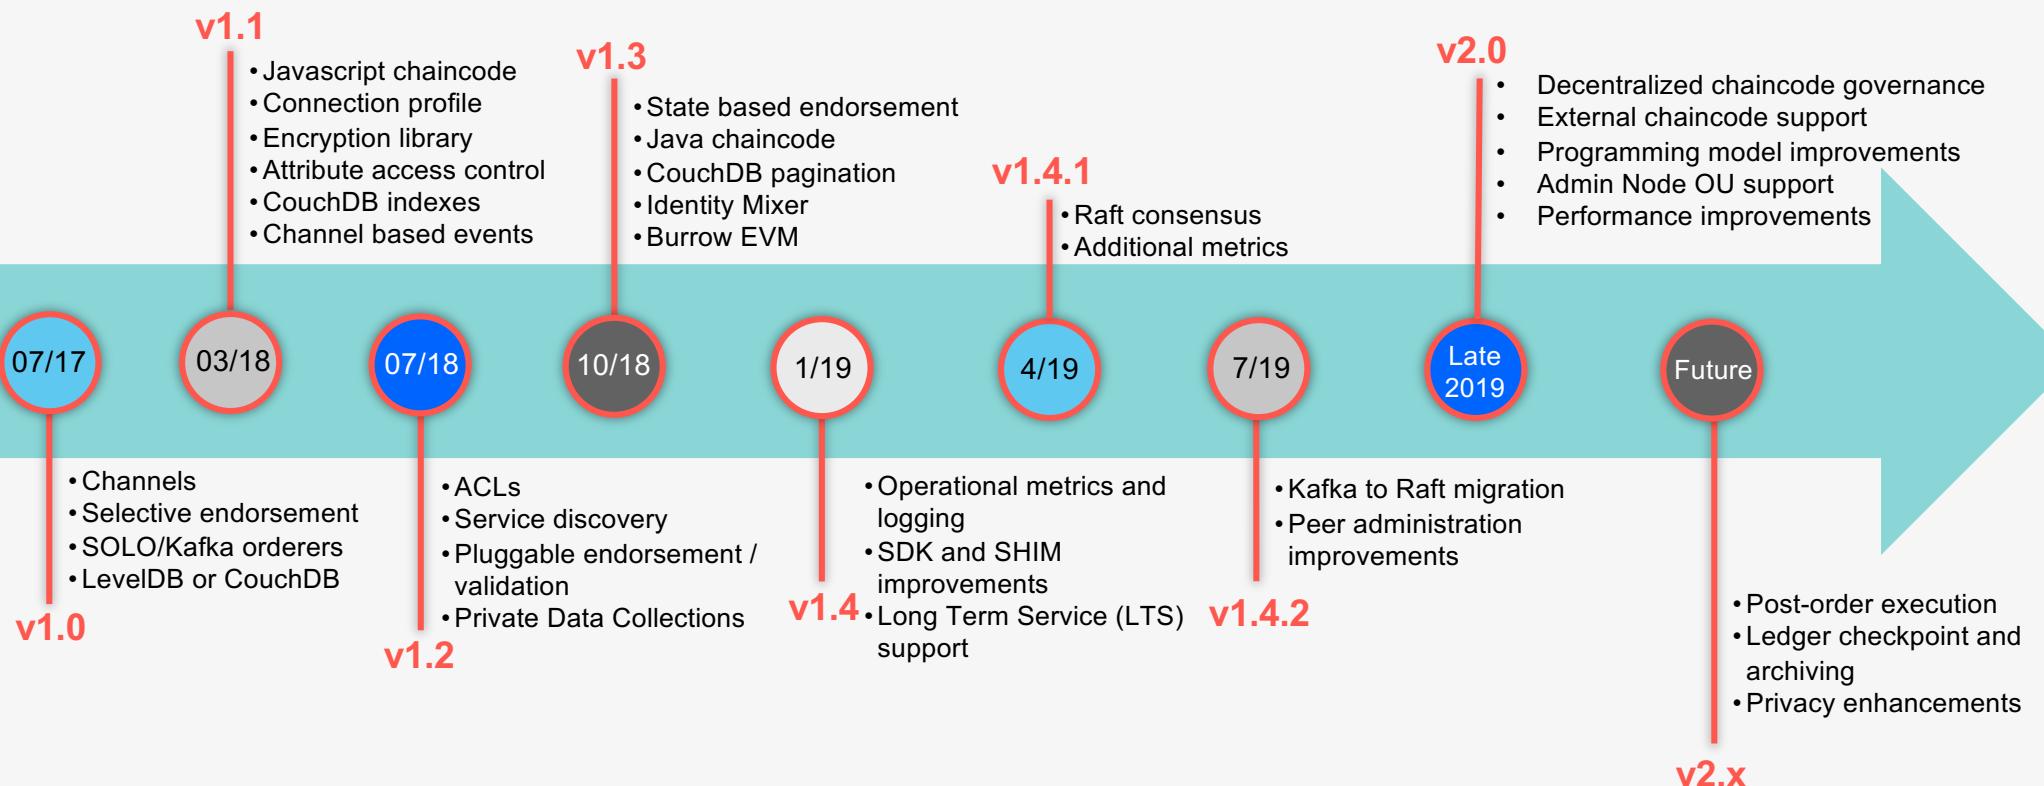
What is Hyperledger Fabric

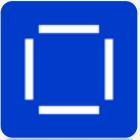


- Linux Foundation Hyperledger
 - A collaborative effort created to advance cross-industry blockchain technologies for business
- Hyperledger Fabric
 - An implementation of blockchain technology that is intended as a foundation for developing blockchain applications
 - Key technical features:
 - A shared ledger and smart contracts implemented as “chaincode”
 - Privacy and permissioning through membership services
 - Modular architecture and flexible hosting options
- First major version released July 2017: contributions by 159 engineers from 27 organizations
 - IBM is one of the contributors to Hyperledger Fabric



Roadmap





IBM Blockchain
Platform

Roadmap

Starter Plan announced

- Entry level developer blockchain environment
- Simple one-click provision and easy simulation of a multi-org environment
- Same look and feel as Enterprise Plan

08/17

02/18

05/18

06/18

10/18

10/18

2Q/19

Starter Plan GA

- Based on Fabric v1.1
- Available in many data centres
- Free 30 day trial
- Provisioned using IBM Containers

Enterprise Plan GA

- Fully managed enterprise grade blockchain-as-a-service
- Built on Fabric 1.0
- Built on LinuxONE

Enterprise Plan v1.1

- New networks provisioned with Fabric v1.1
- Additional data centre locations

Remote Peer Beta announcement

- On-prem peer solution
- Built on Fabric v1.2.1
- Provisioned for IBM Cloud Private

- Next generation platform enabling networks across multi infrastructures
- Developer VSCode plug-in for IBP

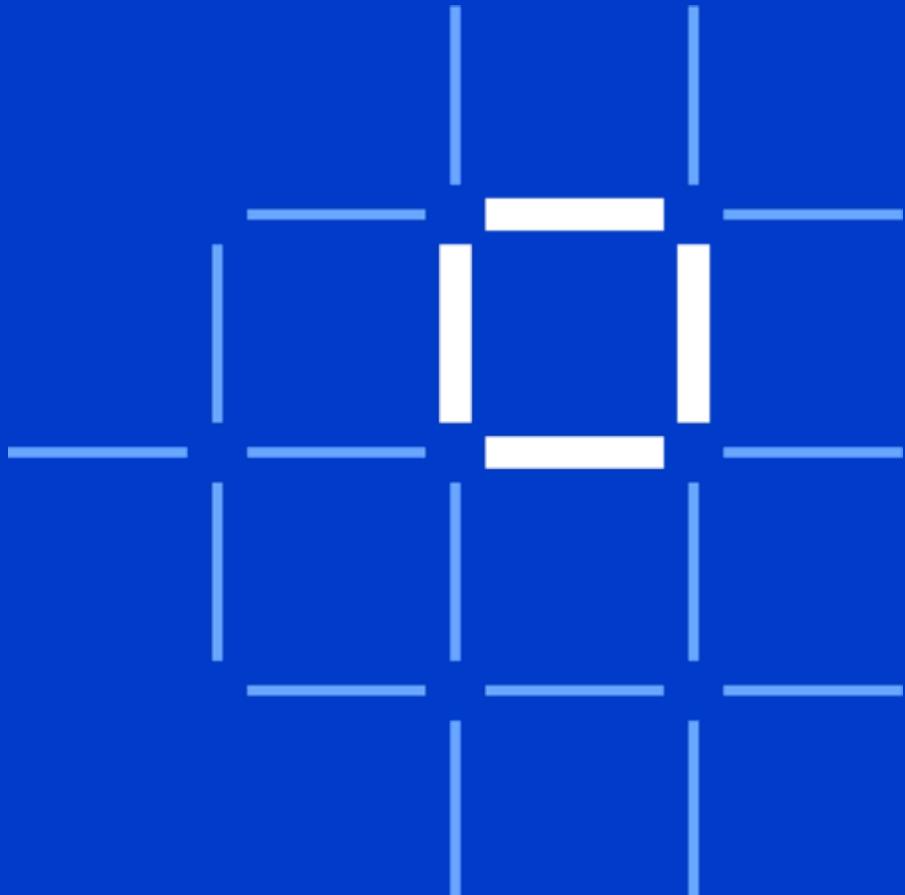
Starter Plan v1.2.1

- Updated to Fabric v1.2.1

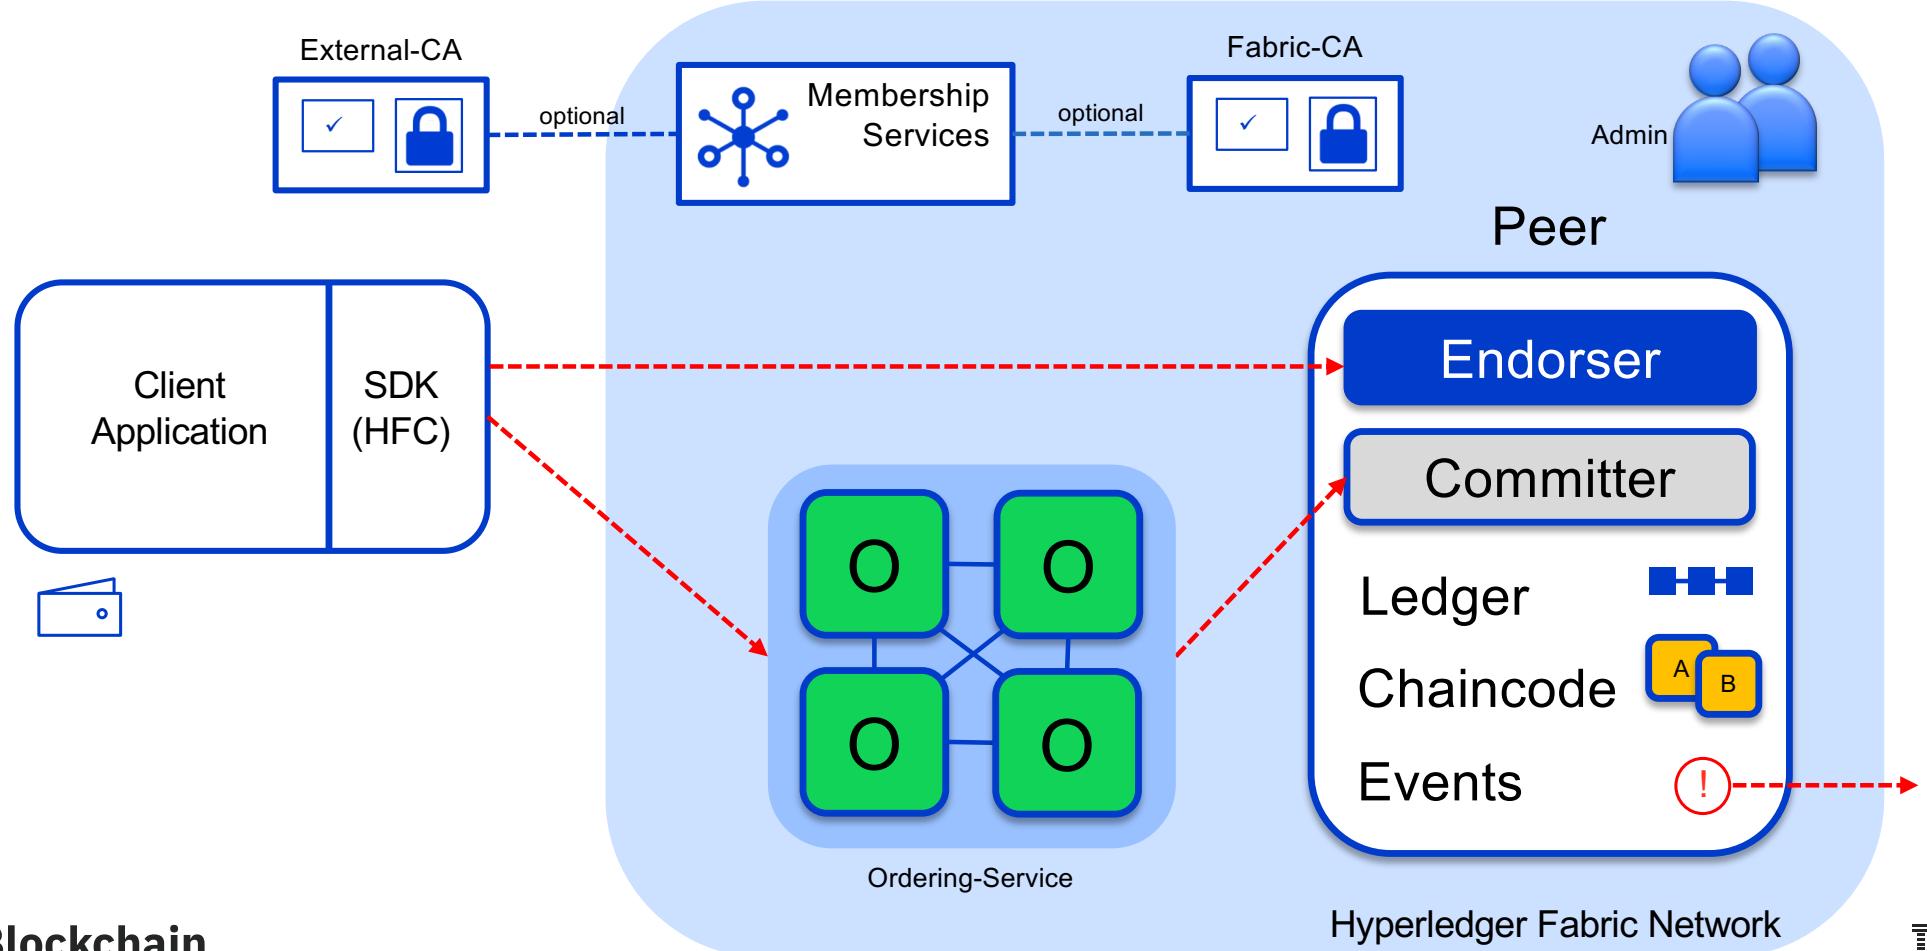
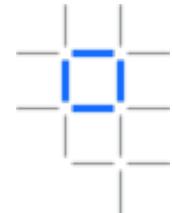


Hyperledger Fabric Technical Dive

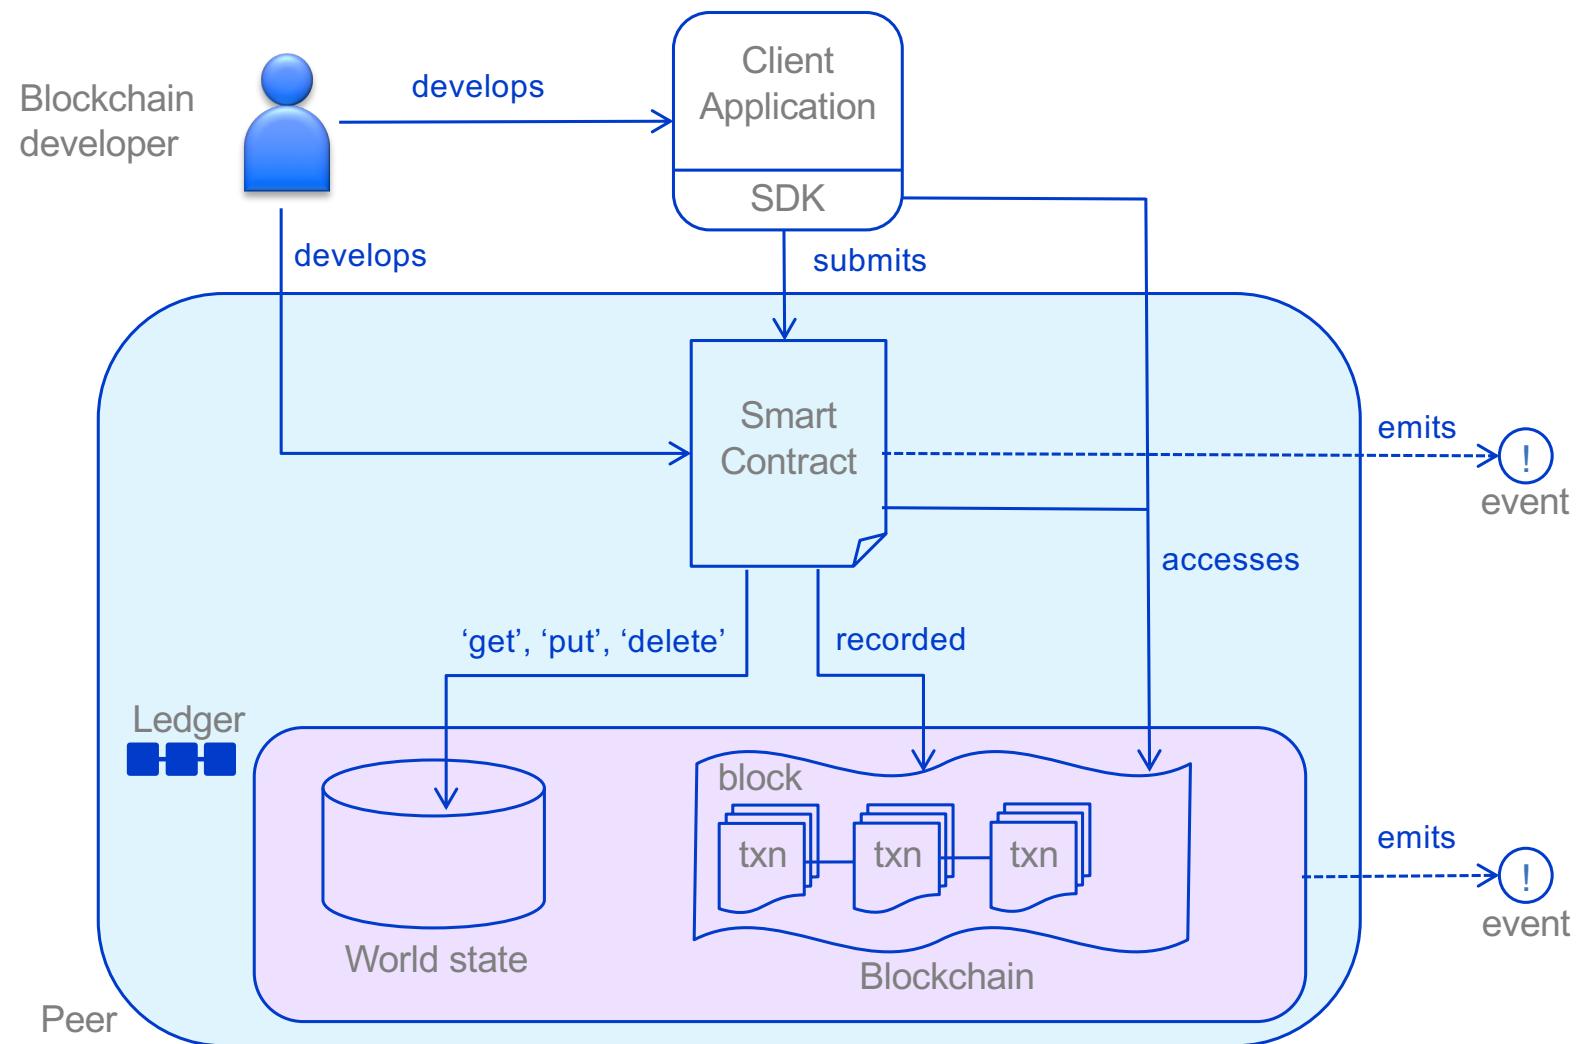
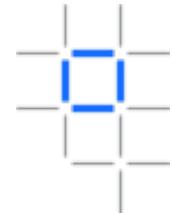
- [Architectural Overview]
- Network Consensus
- Channels and Ordering Service
- Components
- Network setup
- Endorsement Policies
- Membership Services

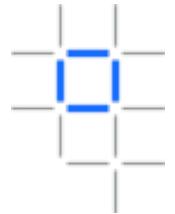


Hyperledger Fabric V1.x Architecture



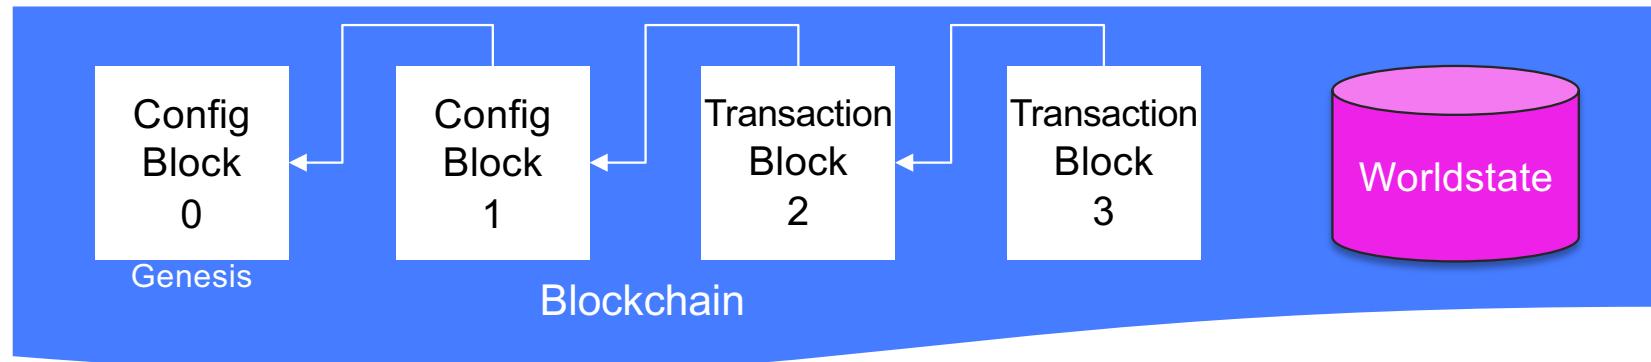
How applications interact with the ledger





Fabric Ledger

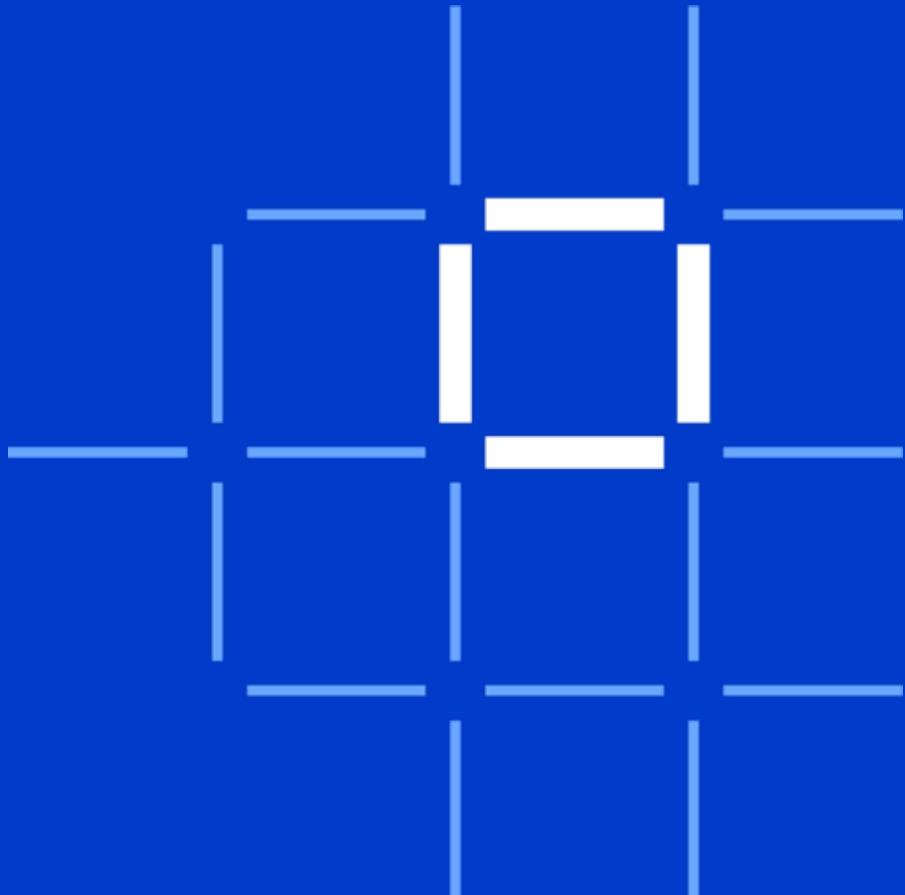
- The **Fabric ledger** is maintained by each peer and includes the **blockchain** and **worldstate**
- A separate ledger is maintained for each channel the peer joins
- Transaction **read/write sets** are written to the blockchain
- **Channel configurations** are also written to the blockchain
- The worldstate can be either LevelDB (default) or CouchDB
 - **LevelDB** is a simple key/value store
 - **CouchDB** is a document store that allows complex queries
- The smart contact Contract decides what is written to the worldstate



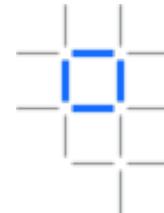


Hyperledger Fabric Technical Dive

- Architectural Overview
- [Network Consensus]
- Channels and Ordering Service
- Components
- Network setup
- Endorsement Policies
- Membership Services

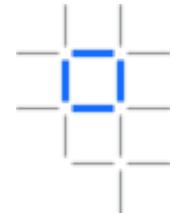


Nodes and roles

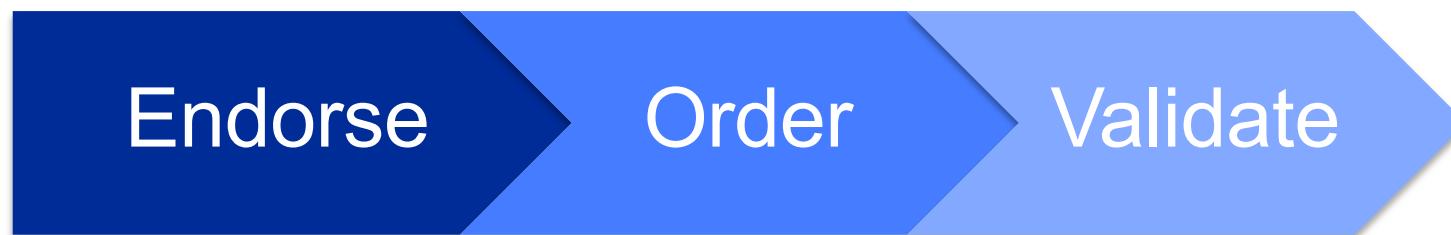


A light gray rounded square with a blue double-line border.	Peer: Maintains ledger and state. Commits transactions. May hold smart contract (chaincode).
A solid blue rounded square with a blue double-line border.	Endorsing Peer: Specialized peer also endorses transactions by receiving a transaction proposal and responds by granting or denying endorsement. Must hold smart contract.
A solid green rounded square with a blue double-line border.	Ordering Node: Approves the inclusion of transaction blocks into the ledger and communicates with committing and endorsing peer nodes. Does not hold smart contract. Does not hold ledger.

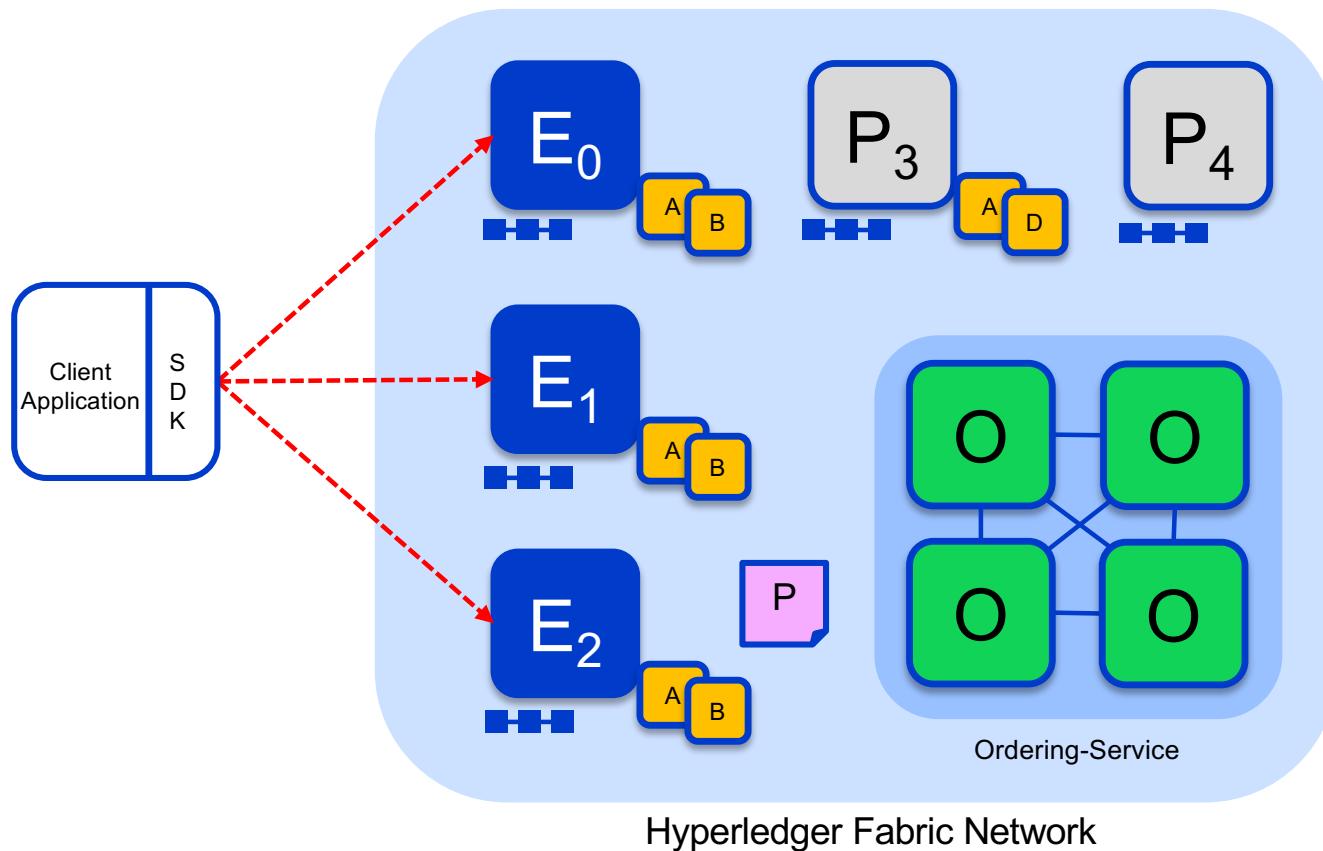
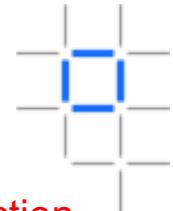
Hyperledger Fabric Consensus



Consensus is achieved using the following transaction flow:



Sample transaction: Step 1/7 – Propose transaction



Application proposes transaction

Endorsement policy:

- " E_0 , E_1 and E_2 must sign"
- (P_3 , P_4 are not part of the policy)

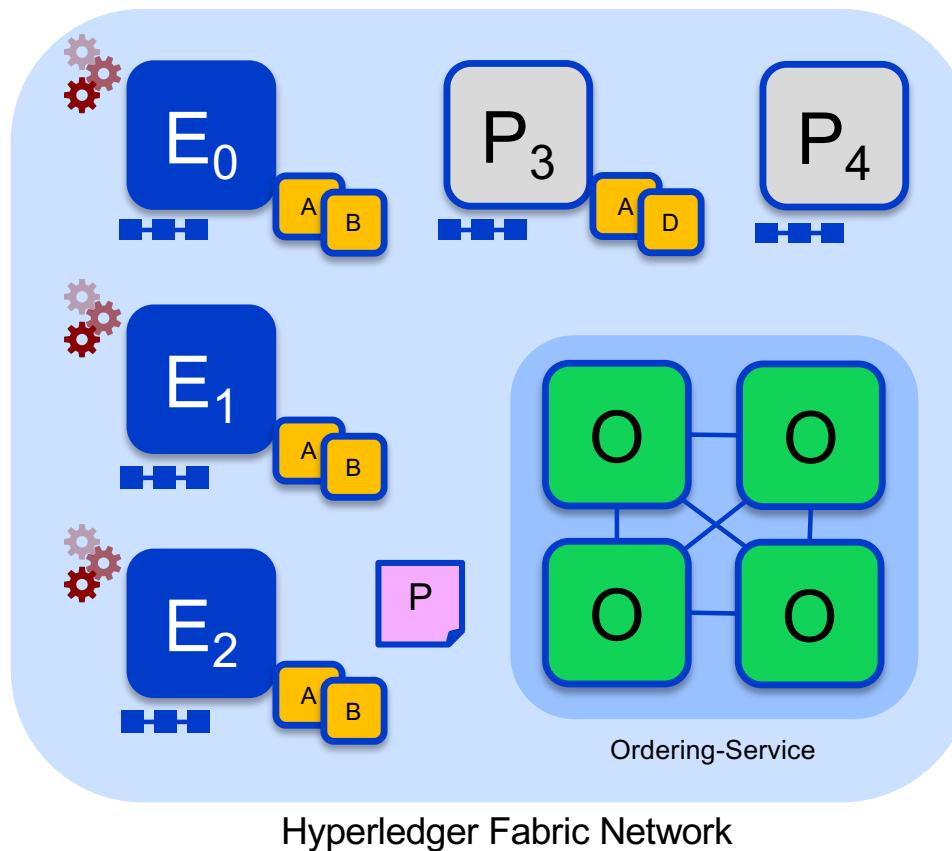
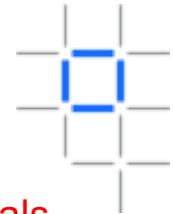
Client application submits a transaction proposal for Smart Contract A. It must target the required peers $\{E_0, E_1, E_2\}$

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy



Sample transaction: Step 2/7 – Execute proposal



Endorsers Execute Proposals

E_0 , E_1 & E_2 will each execute the proposed transaction. None of these executions will update the ledger

Each execution will capture the set of Read and Written data, called RW sets, which will now flow in the fabric.

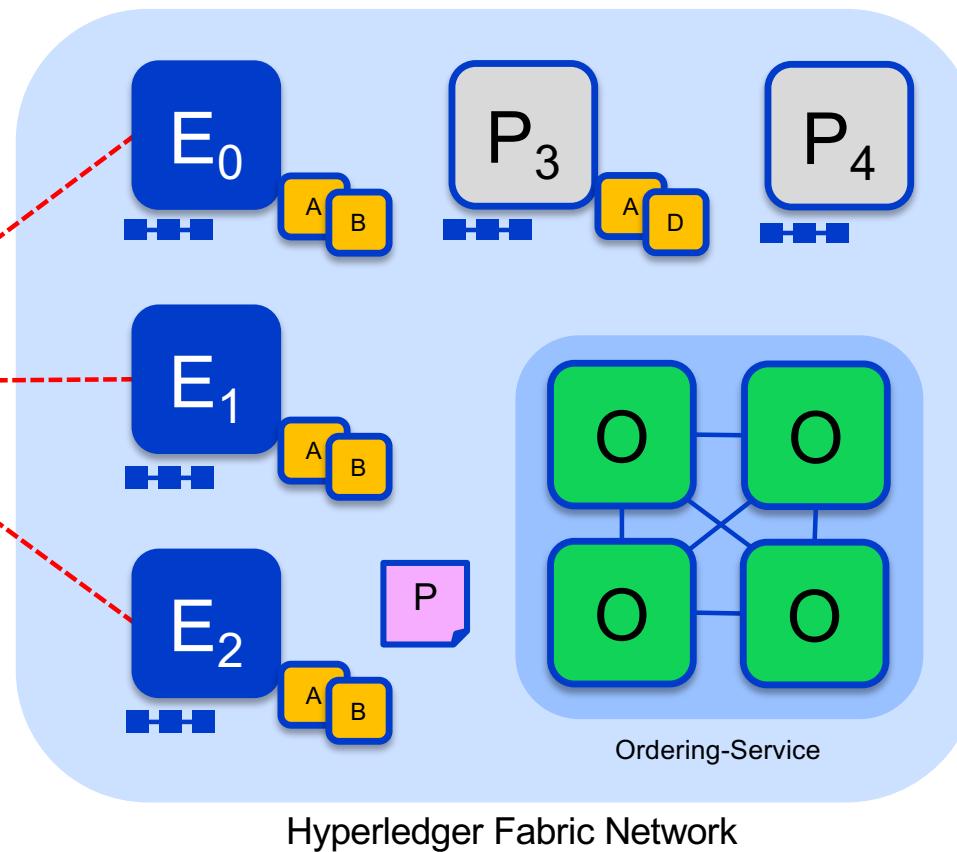
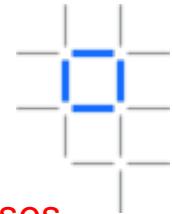
Transactions can be signed & encrypted

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy



Sample transaction: Step 3/7 – Proposal Response



Application receives responses

RW sets are asynchronously returned to application

The RW sets are signed by each endorser, and also includes each record version number

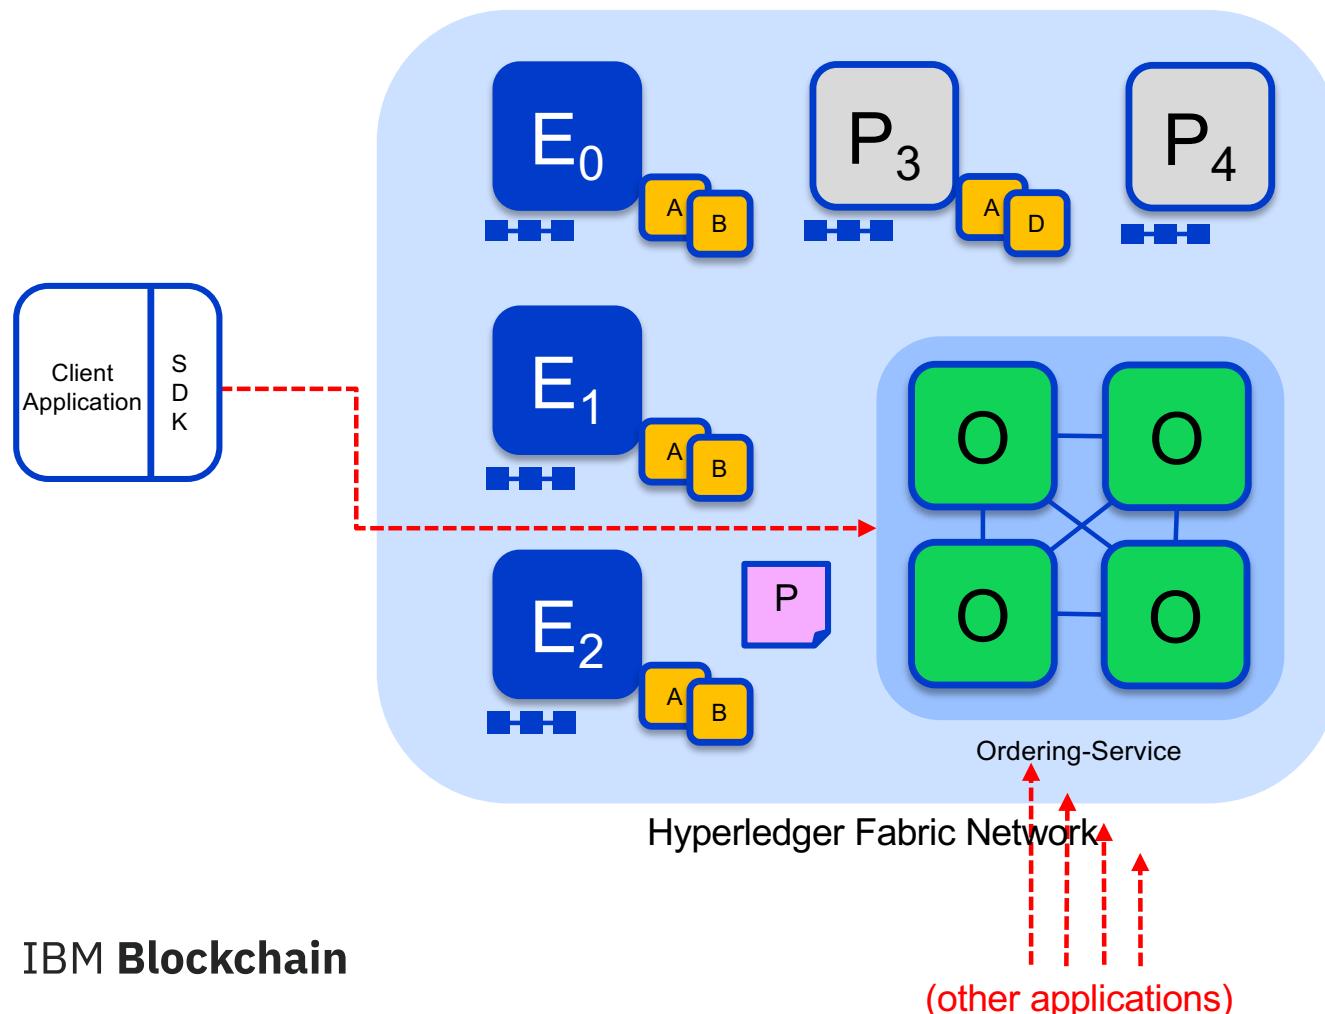
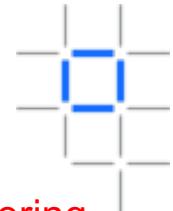
(This information will be checked much later in the consensus process)

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy



Sample transaction: Step 4/7 – Order Transaction



Responses submitted for ordering

Application submits responses as a transaction to be ordered.

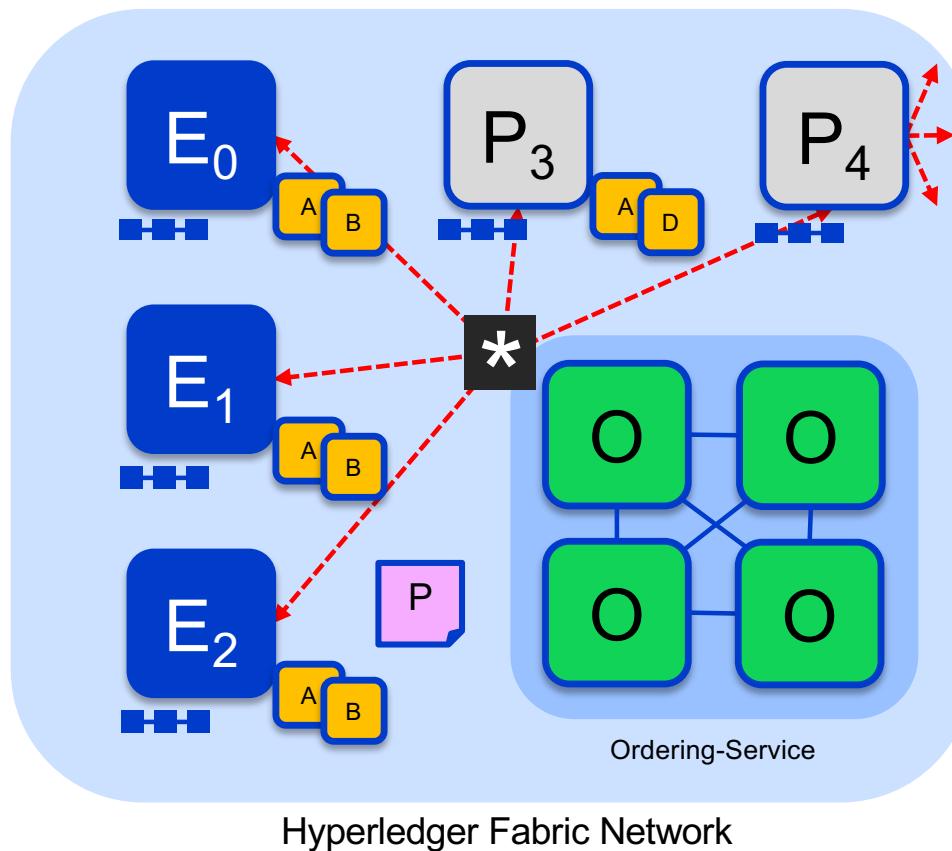
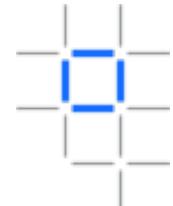
Ordering happens across the fabric in parallel with transactions submitted by other applications

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy



Sample transaction: Step 5/7 – Deliver Transaction



Orderer delivers to committing peers

Ordering service collects transactions into proposed blocks for distribution to committing peers. Peers can deliver to other peers in a hierarchy (not shown)

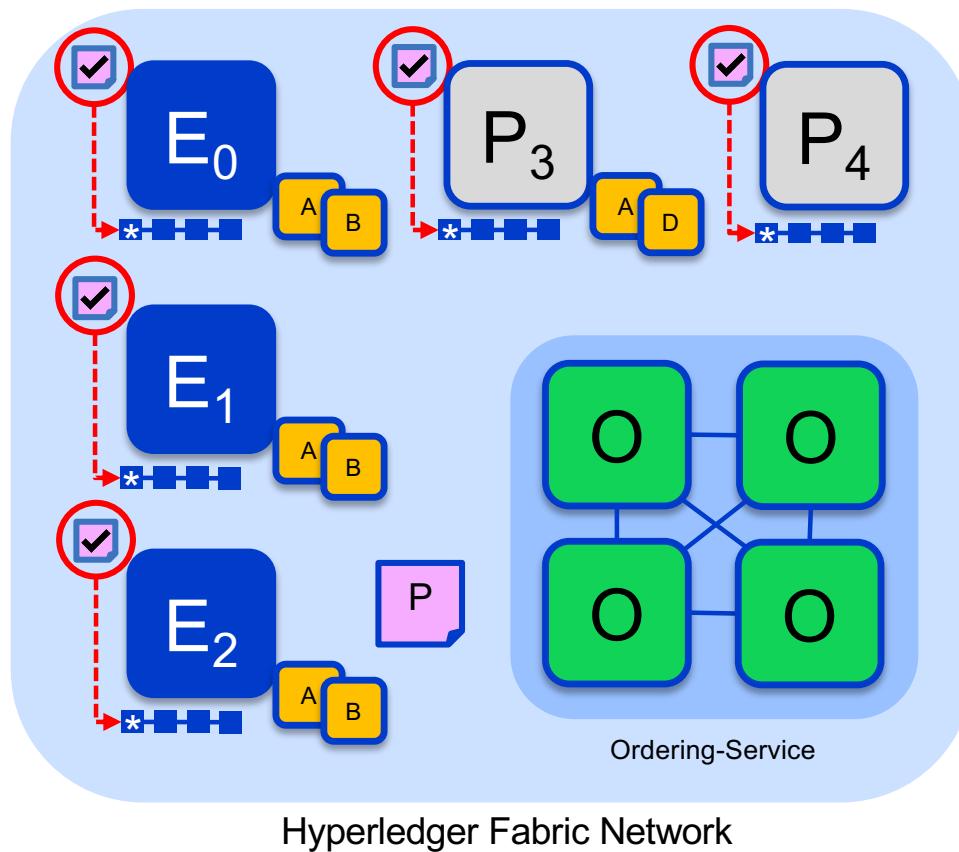
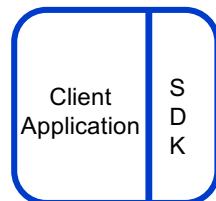
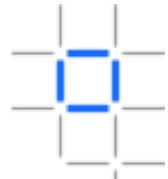
RAFT consensus provides crash fault tolerance

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy



Sample transaction: Step 6/7 – Validate Transaction



Committing peers validate transactions

Every committing peer validates against the endorsement policy. Also check RW sets are still valid for current world state

Validated transactions are applied to the world state and retained on the ledger

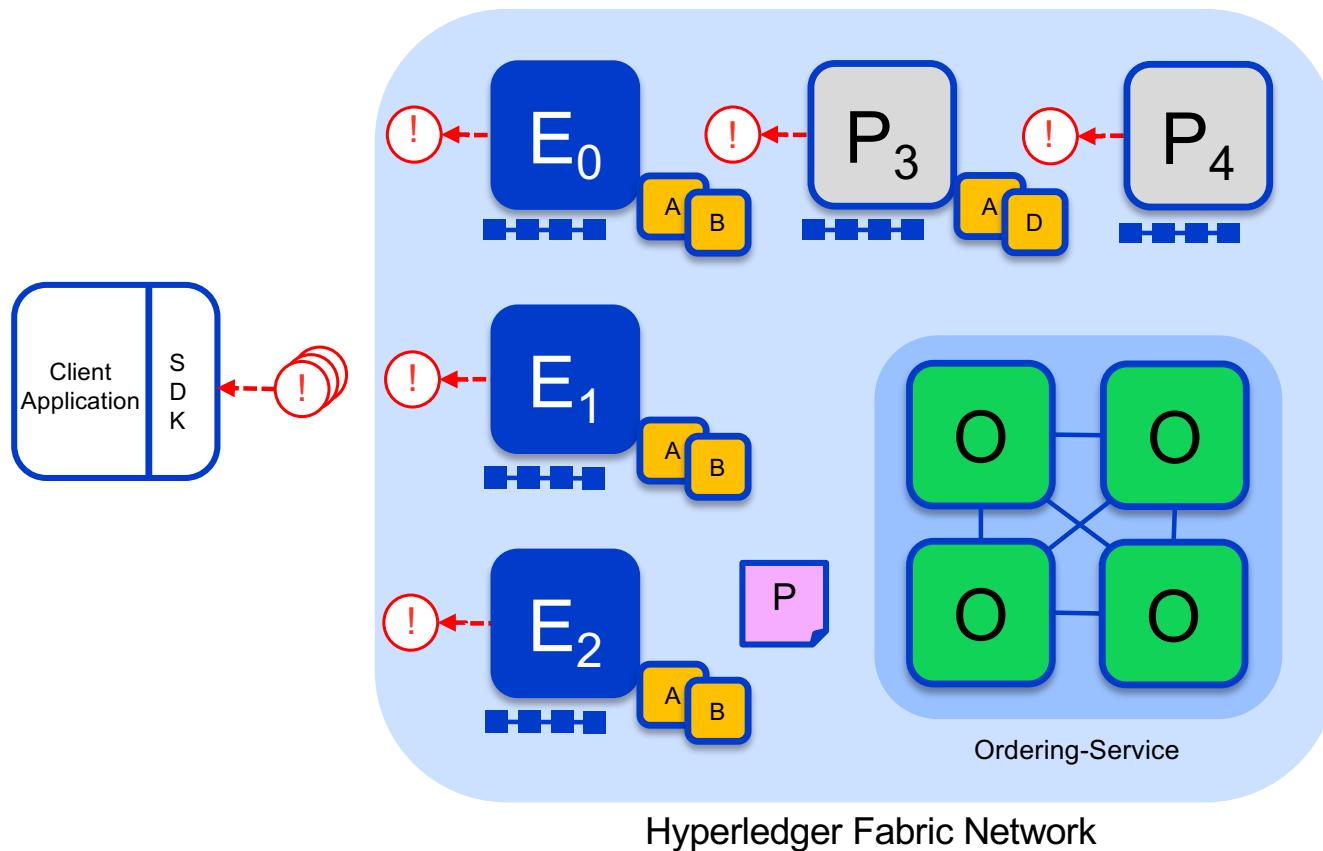
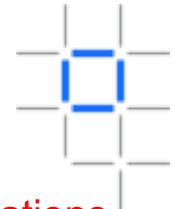
Invalid transactions are also retained on the ledger but do not update world state

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy



Sample transaction: Step 7/7 – Notify Transaction



Committing peers notify applications

Applications can register to be notified when transactions succeed or fail, and when blocks are added to the ledger

Applications will be notified by each peer to which they are connected

Key:

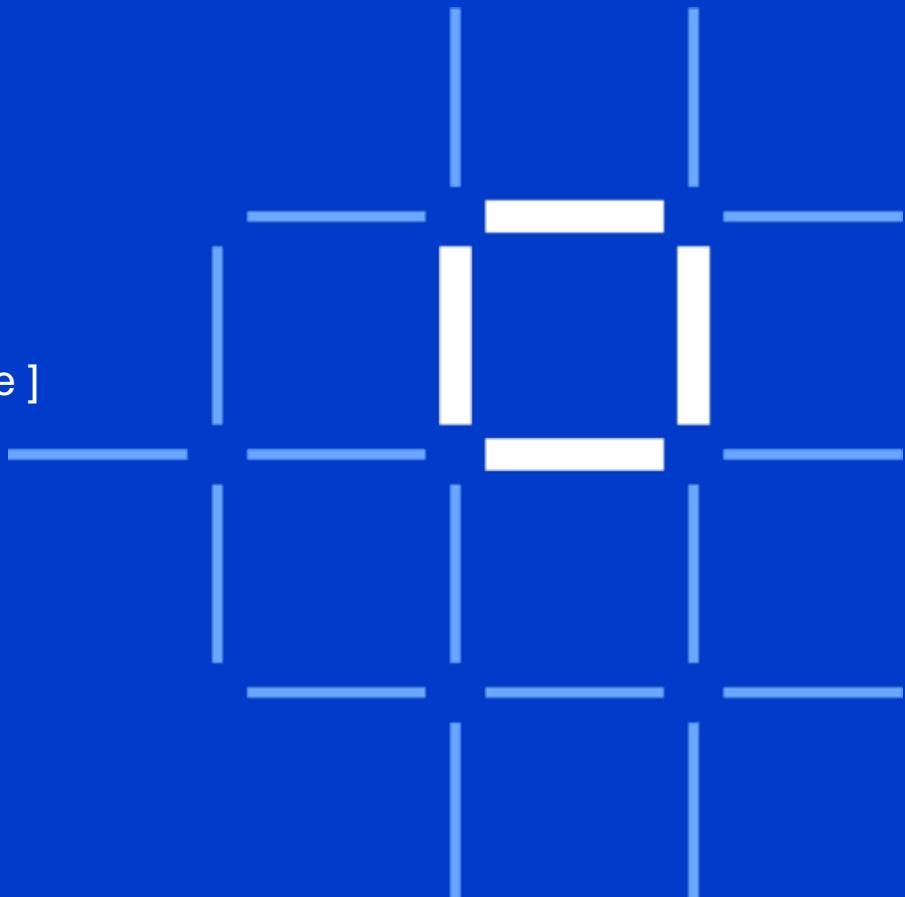
Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy



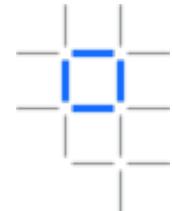


Hyperledger Fabric Technical Dive

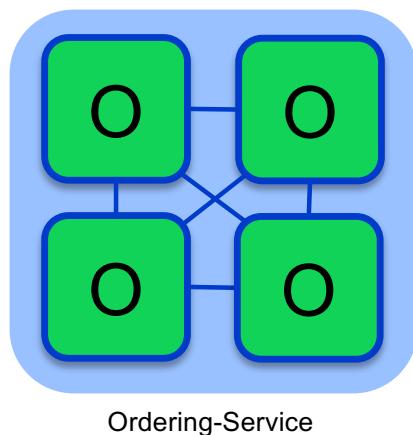
- Architectural Overview
- Network Consensus
- [Channels and Ordering Service]
- Components
- Network setup
- Endorsement Policies
- Membership Services



Ordering Service



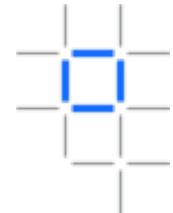
The ordering service packages transactions into blocks to be delivered to peers. Communication with the service is via channels.



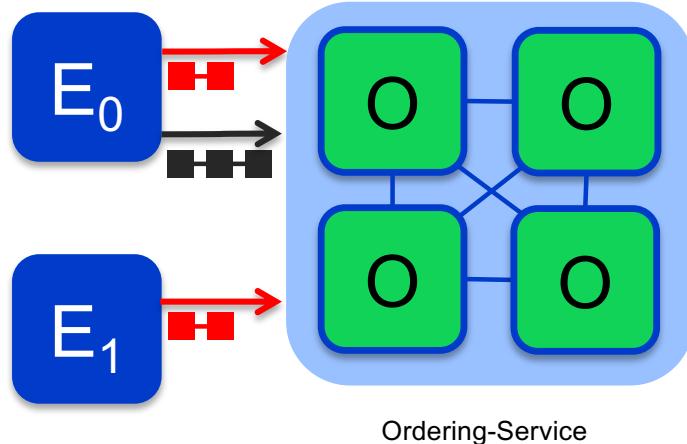
- RAFT consensus for the ordering service

- Crash fault tolerant consensus
- 3 nodes minimum
- Odd number of nodes recommended

Channels

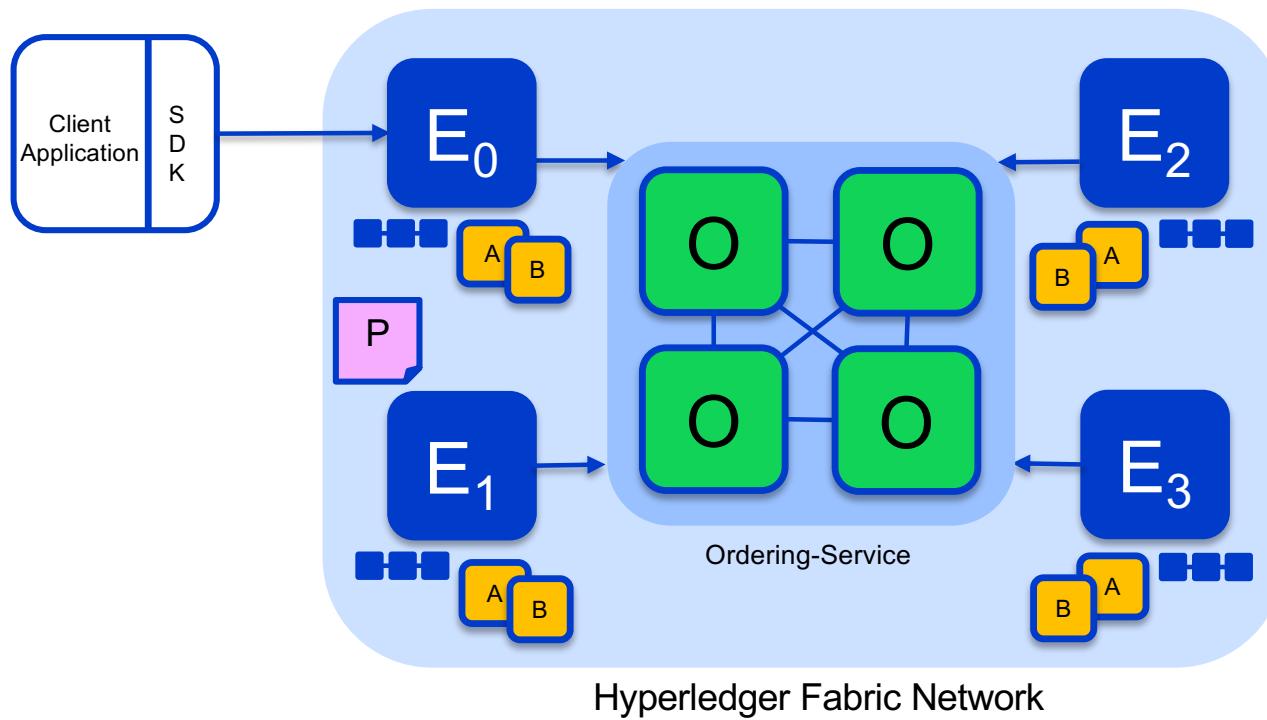
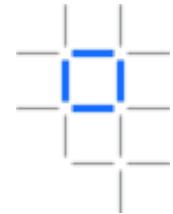


Channels provide privacy between different ledgers



- Ledgers exist in the scope of a channel
 - Channels can be shared across an entire network of peers
 - Channels can be permissioned for a specific set of participants
- Chaincode is **installed** on peers to access the worldstate
- Chaincode is **instantiated** on specific channels
- Peers can participate in multiple channels
- Concurrent execution for performance and scalability

Single Channel Network

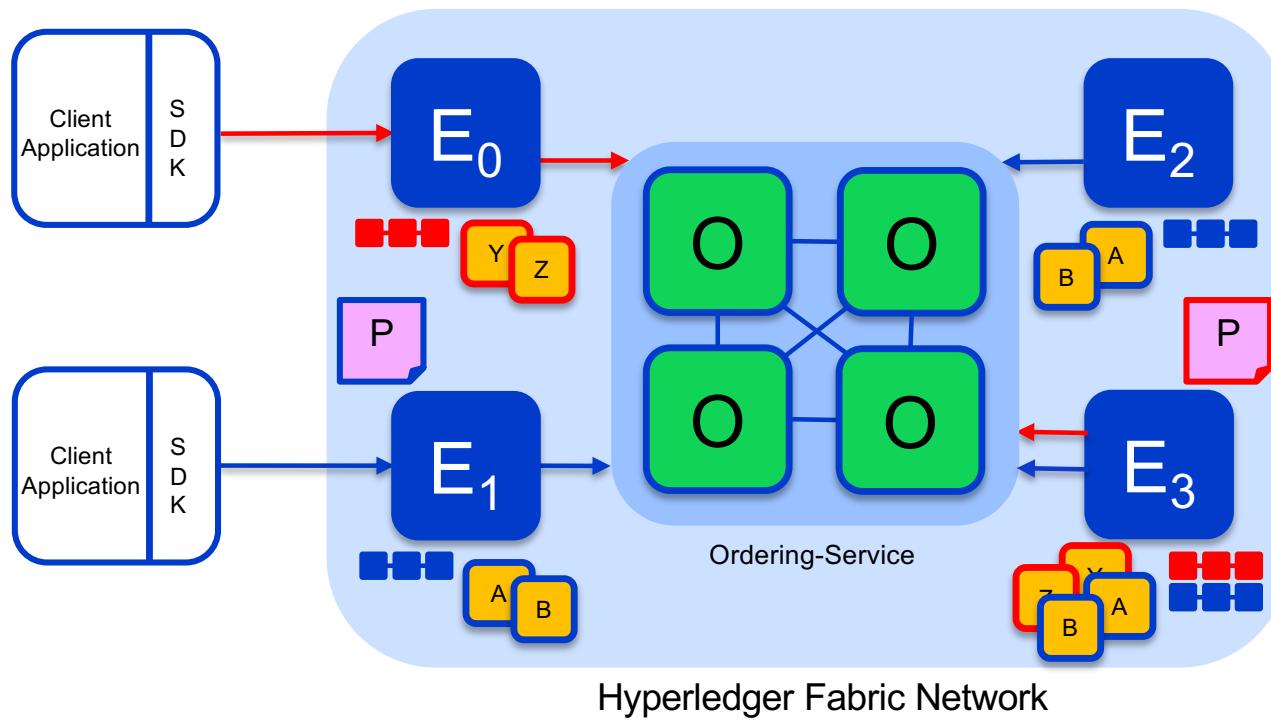
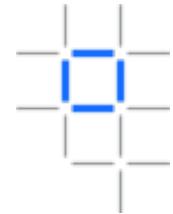


- All peers connect to the same system channel (blue).
- All peers have the same chaincode and maintain the same ledger
- Endorsement by peers E_0 , E_1 , E_2 and E_3

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy

Multi Channel Network



- Peers E_0 and E_3 connect to the red channel for chaincodes Y and Z
- E_1 , E_2 and E_3 connect to the blue channel for chaincodes A and B

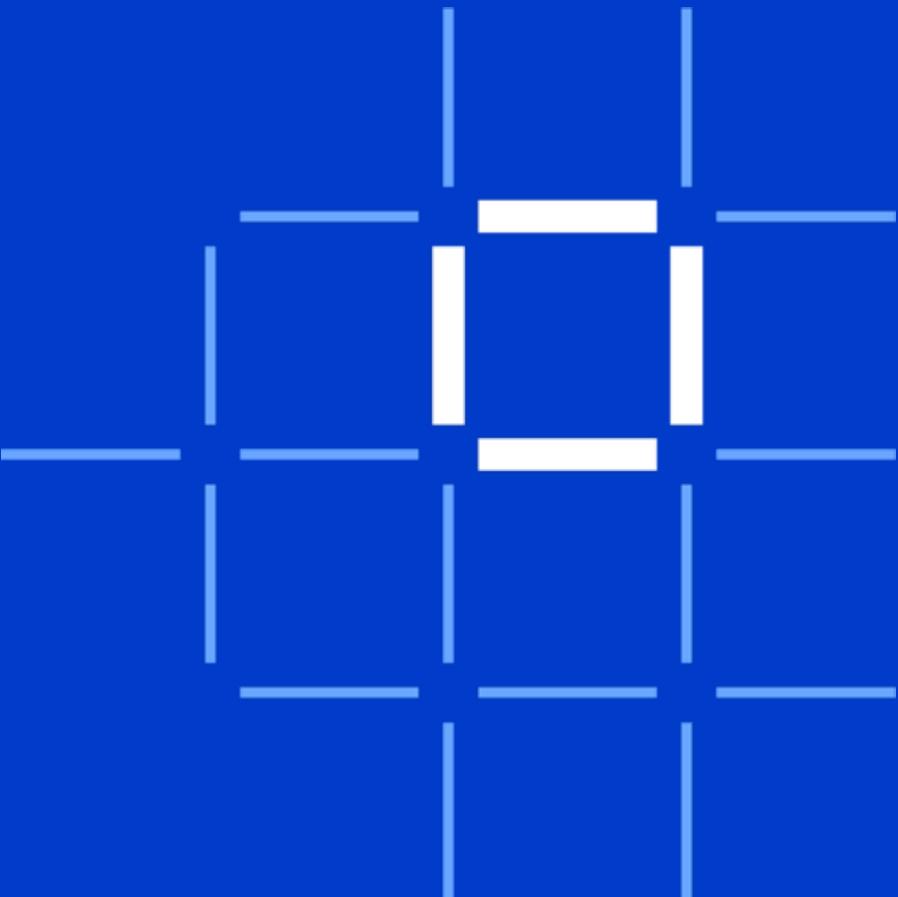
Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy



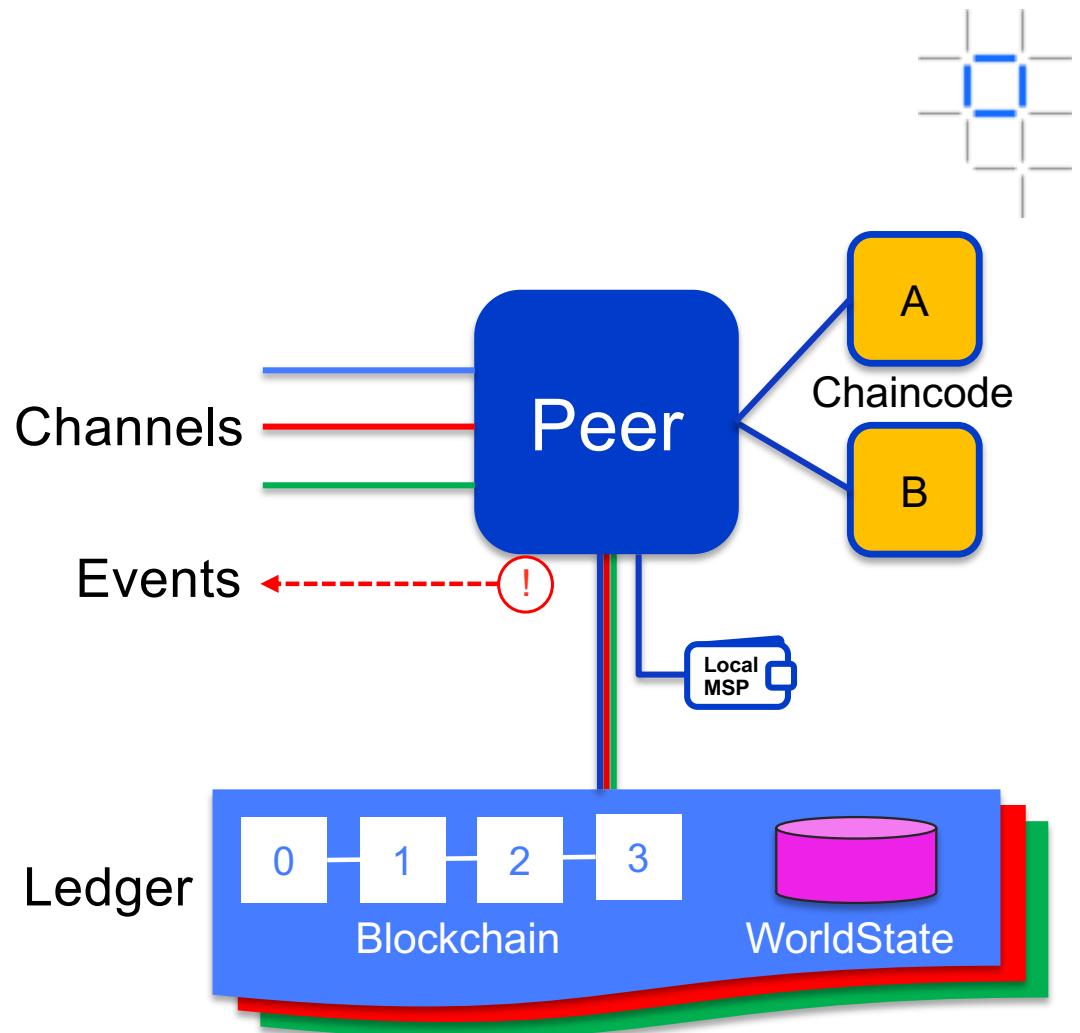
Hyperledger Fabric Technical Dive

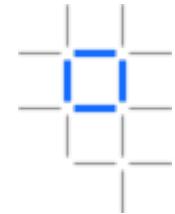
- Architectural Overview
- Network Consensus
- Channels and Ordering Service
- [Components]
- Network setup
- Endorsement Policies
- Membership Services



Fabric Peer

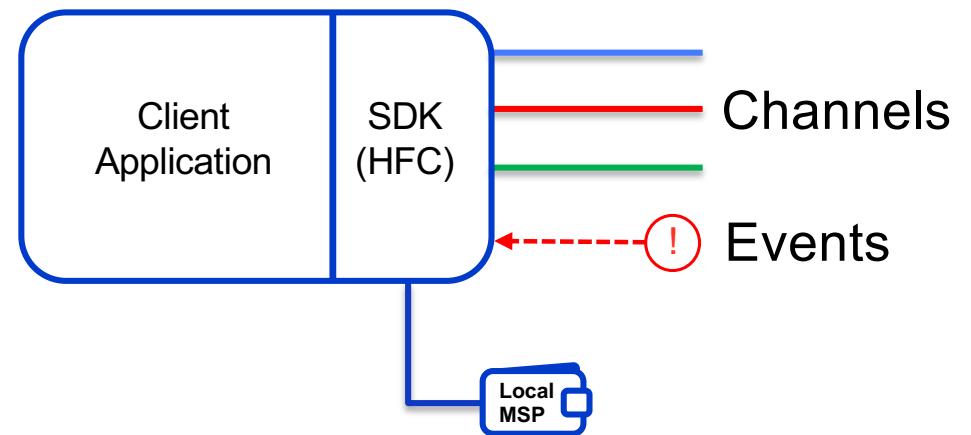
- Each peer:
 - Connects to one or more **channels**
 - Maintains one or more **ledgers** per channel
 - Maintains **installed chaincode**
 - Manages **runtime docker containers** for **instantiated chaincode**
 - Chaincode is instantiated on a channel
 - Runtime docker container shared by channels with same chaincode instantiated (no state stored in container)
 - Has a local MSP (Membership Services Provider) that provides **crypto material**
 - **Emits events** to the client application

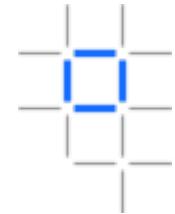




Client Application

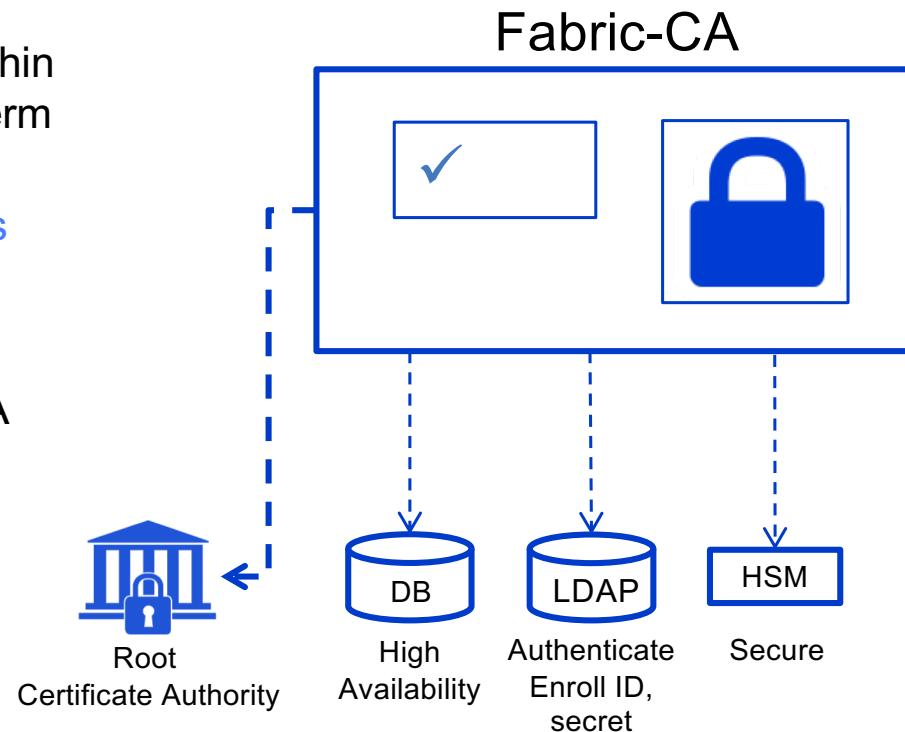
- Each client application uses Fabric SDK to:
 - Connects over channels to one or more peers
 - Connects over channels to one or more orderer nodes
 - Receives events from peers
 - Local MSP provides client [crypto material](#)
- Client can be written in different languages
(Node.js, Go, Java, Python?)





Fabric-CA

- Default (optional) Certificate Authority within Fabric network for issuing **Ecerts** (long-term identity)
- Supports clustering for **HA characteristics**
- Supports LDAP for **user authentication**
- Supports HSM for **security**
- Can be configured as an intermediate CA



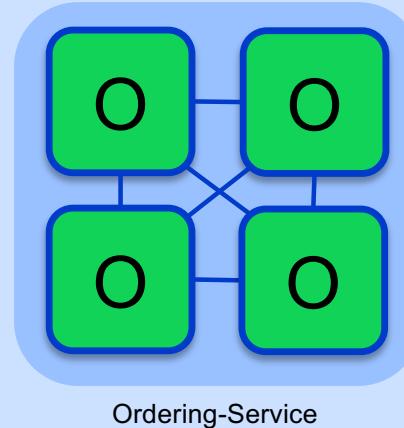
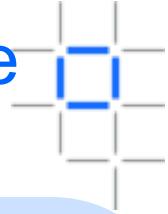


Hyperledger Fabric Technical Dive

- Architectural Overview
- Network Consensus
- Channels and Ordering Service
- Components
- [Network setup]
- Endorsement Policies
- Membership Services



Bootstrap Network (1/6) - Configure & Start Ordering Service

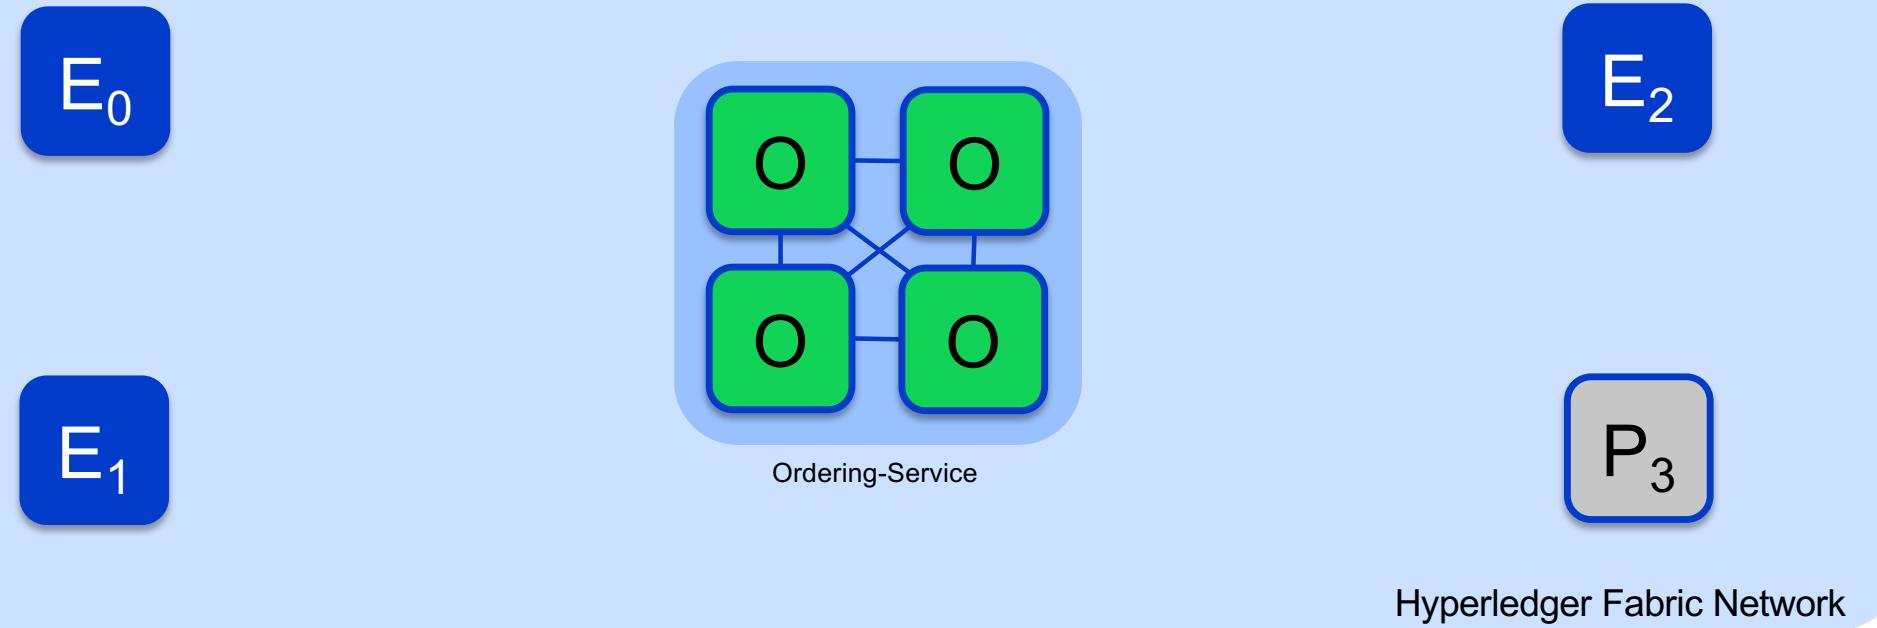
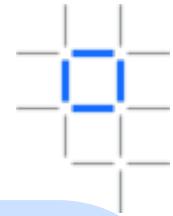


Hyperledger Fabric Network

An Ordering Service is configured and started for the network:

\$ docker-compose [-f orderer.yml] ...

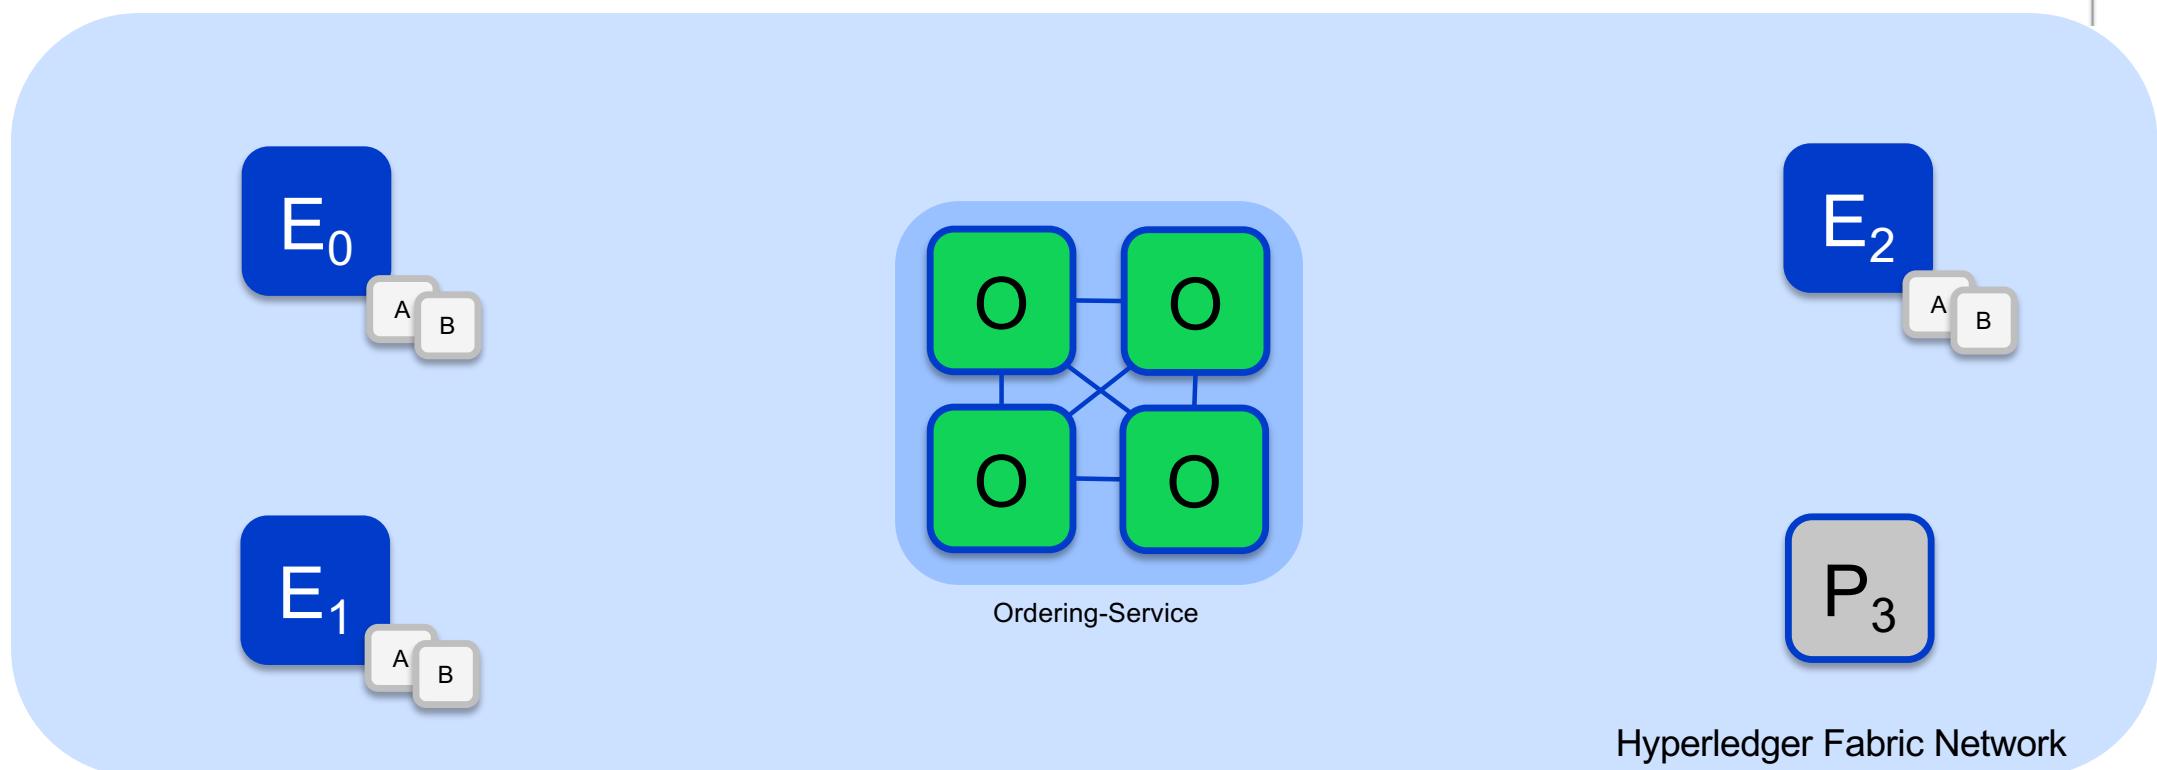
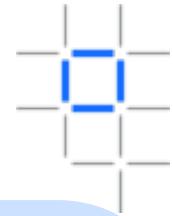
Bootstrap Network (2/6) - Configure and Start Peer Nodes



A peer is configured and started for each Endorser or Committer in the network:

\$ peer node start ...

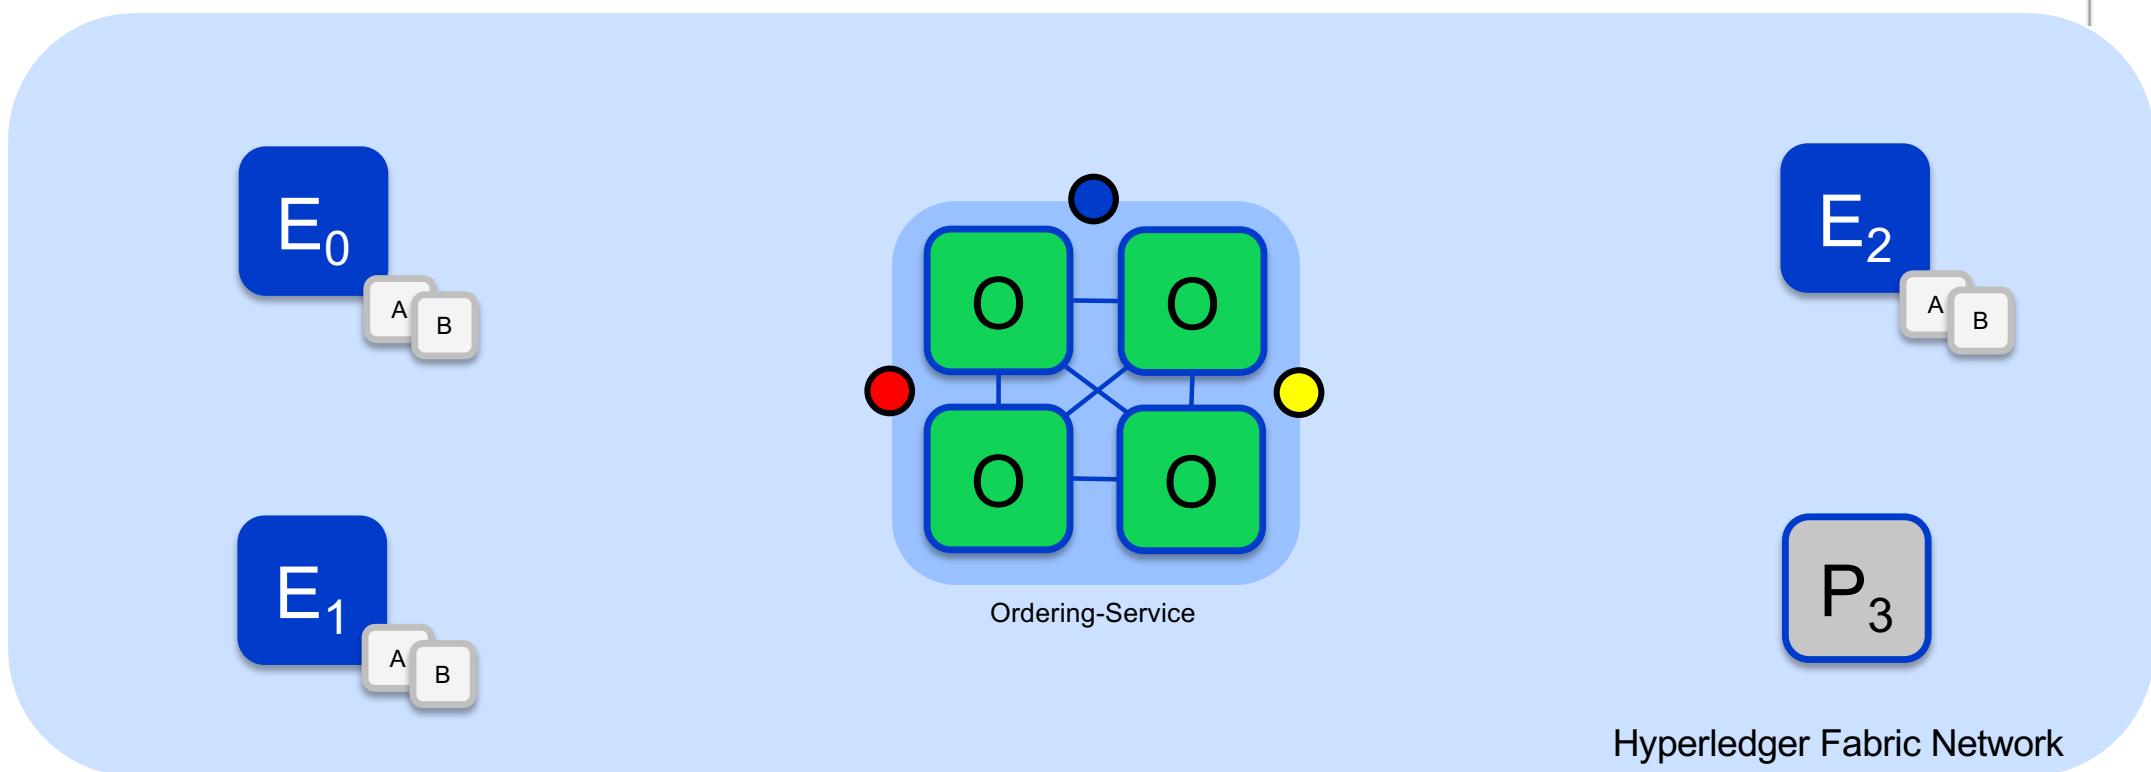
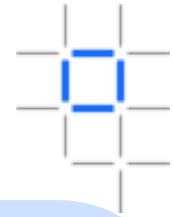
Bootstrap Network (3/6) - Install Chaincode



Chaincode is installed onto each Endorsing Peer that needs to execute it:

\$ peer chaincode install ...

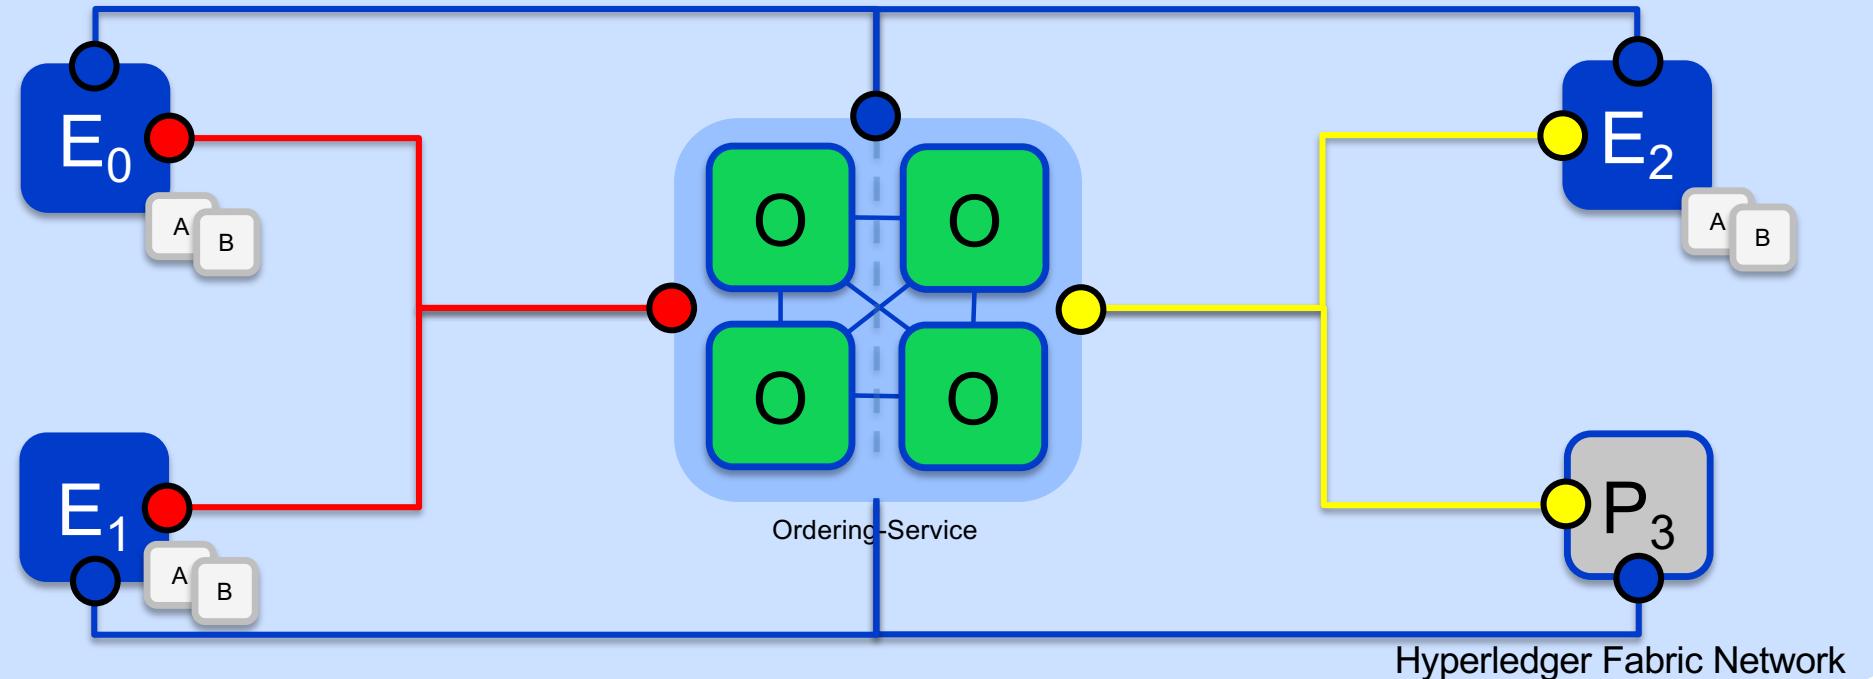
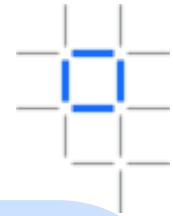
Bootstrap Network (4/6) – Create Channels



Channels are created on the ordering service:

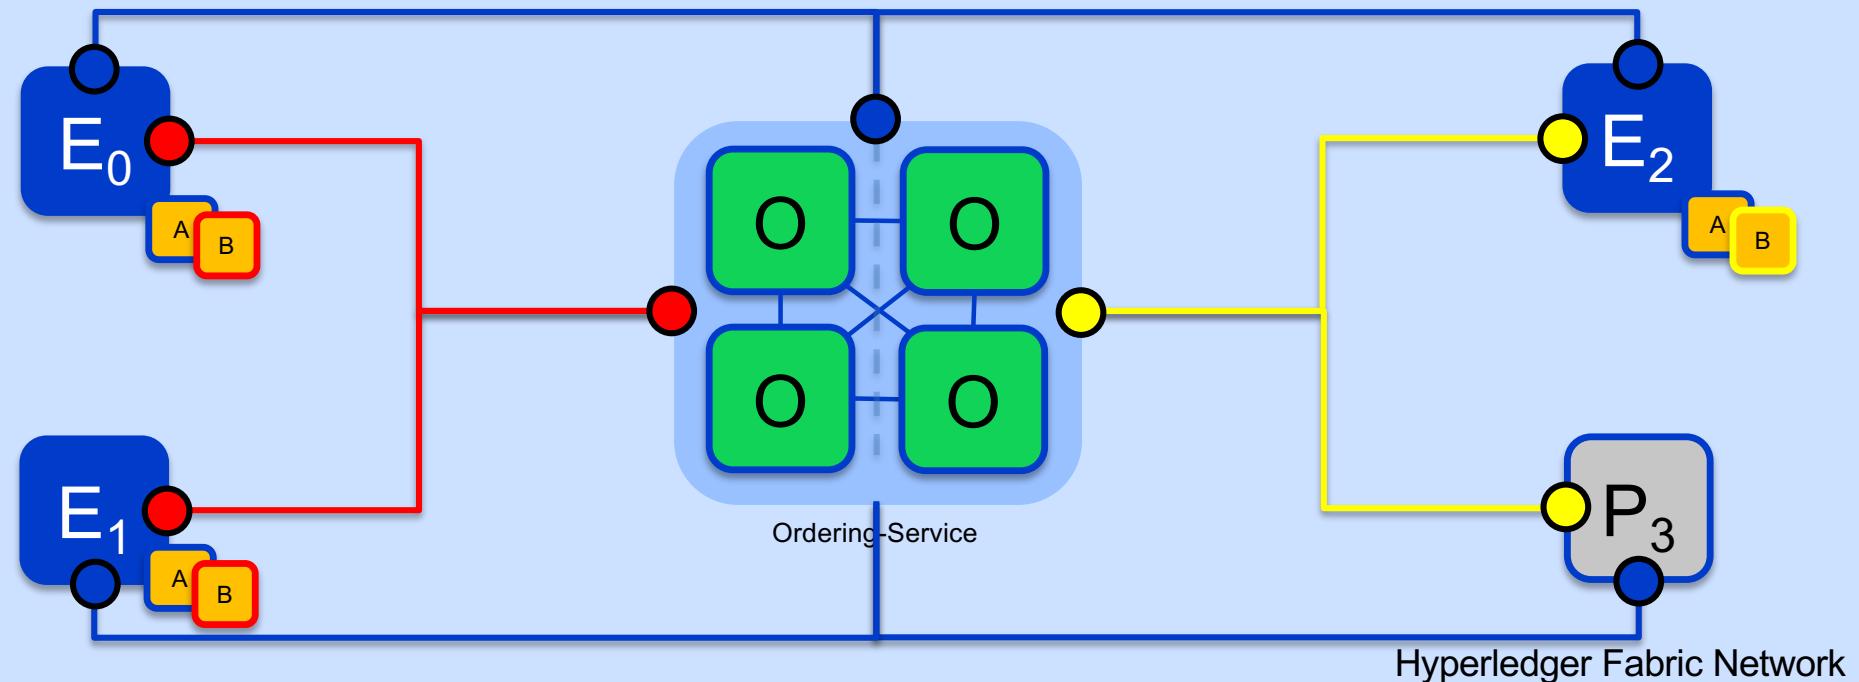
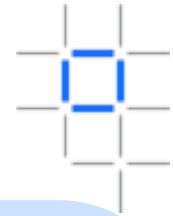
```
$ peer channel create -o [orderer] ...
```

Bootstrap Network (5/6) – Join Channels



Peers that are permissioned can then join the channels they want to transact on:
\$ peer channel join ...

Bootstrap Network (6/6) – Instantiate Chaincode

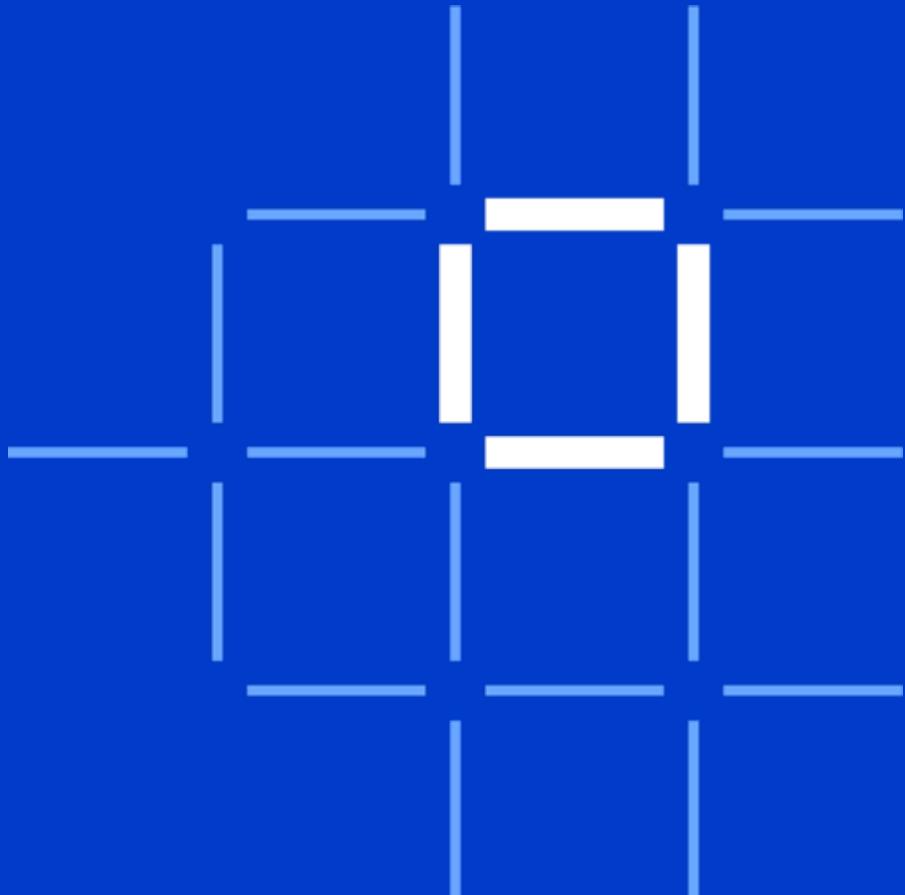


Peers finally instantiate the Chaincode on the channels they want to transact on:
\$ peer chaincode instantiate ... –P ‘policy’

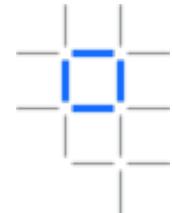


Hyperledger Fabric Technical Dive

- Architectural Overview
- Network Consensus
- Channels and Ordering Service
- Components
- Network setup
- [Endorsement Policies]
- Membership Services

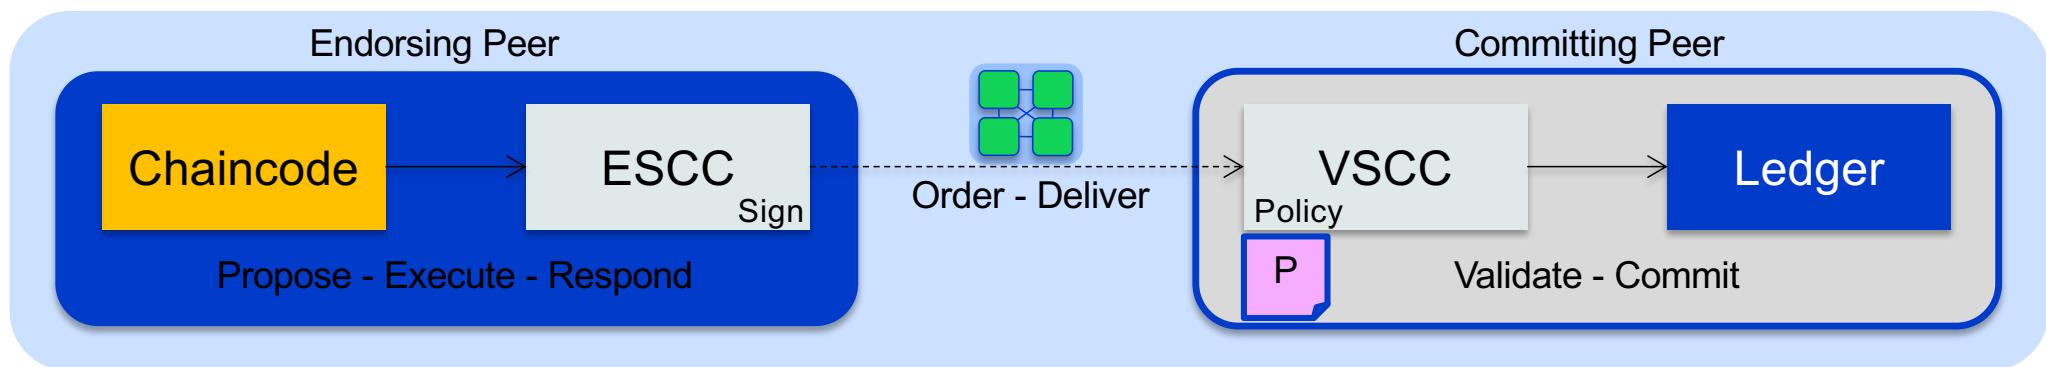


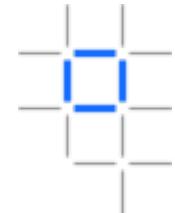
Endorsement Policies



An endorsement policy describes the conditions by which a transaction can be endorsed. A transaction can only be considered valid if it has been endorsed according to its policy.

- Each chaincode is deployed with an Endorsement Policy
- **ESCC** ([Endorsement System ChainCode](#)) signs the proposal response on the endorsing peer
- **VSCC** ([Validation System ChainCode](#)) validates the endorsements





Endorsement Policy Syntax

```
$ peer chaincode instantiate  
-C mychannel  
-n mycc  
-v 1.0  
-p chaincode_example02  
-c '{"Args":["init","a", "100", "b", "200"] }'  
-P "AND ('Org1MSP.member') "
```

Instantiate the chaincode **mycc** on channel **mychannel** with the policy **AND('Org1MSP.member')**

Policy Syntax: **EXPR(E[, E...])**

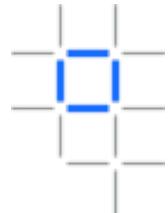
Where **EXPR** is either AND or OR and **E** is either a principal or nested EXPR

Principal Syntax: **MSP.ROLE**

Supported roles are: member and admin

Where **MSP** is the MSP ID, and **ROLE** is either “member” or “admin”

Endorsement Policy Examples



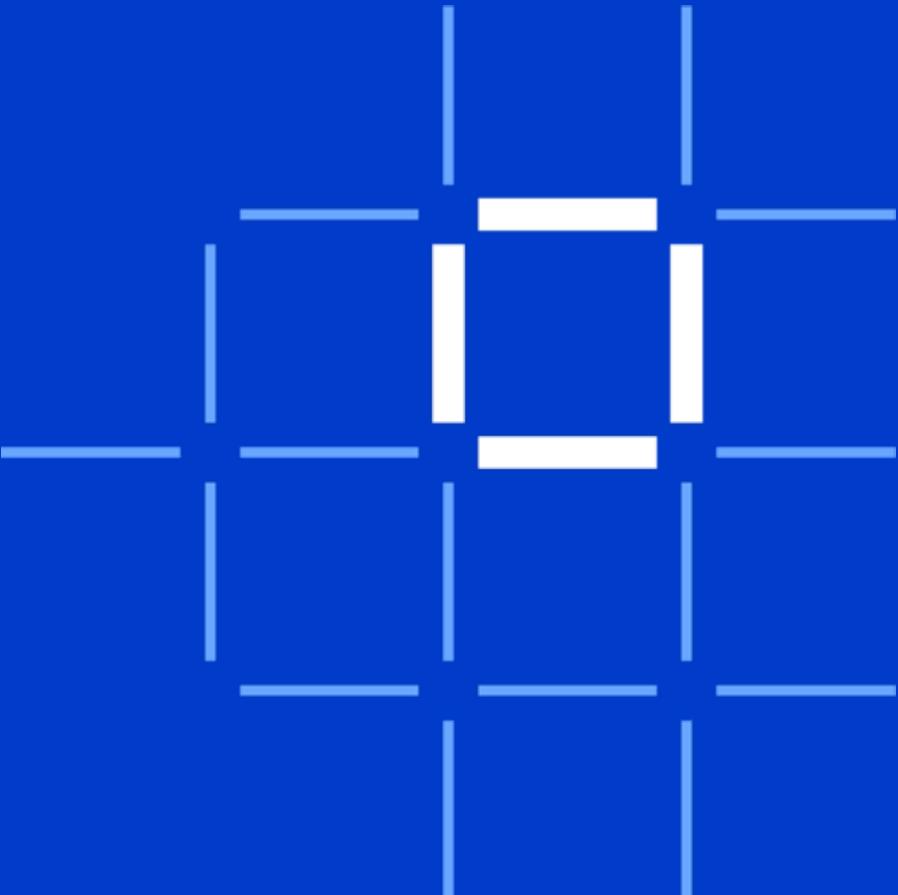
Examples of policies:

- Request 1 signature from all three principals
 - `AND('Org1.member', 'Org2.member', 'Org3.member')`
- Request 1 signature from either one of the two principals
 - `OR('Org1.member', 'Org2.member')`
- Request either one signature from a member of the Org1 MSP or (1 signature from a member of the Org2 MSP and 1 signature from a member of the Org3 MSP)
 - `OR('Org1.member', AND('Org2.member', 'Org3.member'))`

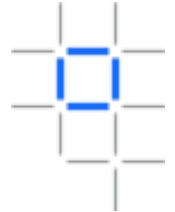


Hyperledger Fabric Technical Dive

- Architectural Overview
- Network Consensus
- Channels and Ordering Service
- Components
- Network setup
- Endorsement Policies
- [Membership Services]

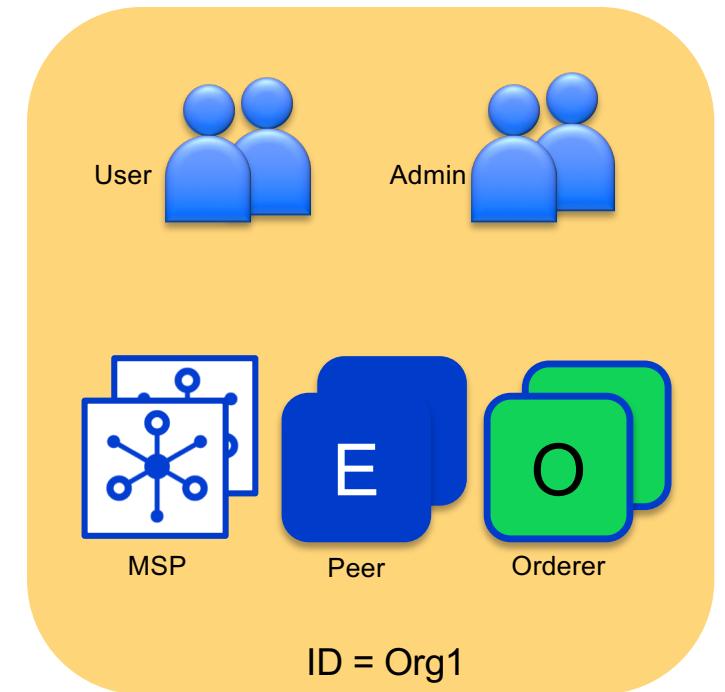


Organisations

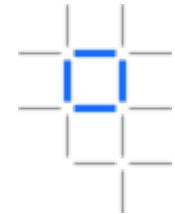


Organisations define boundaries within a Fabric Blockchain Network

- Each organisation defines:
 - Membership Services Provider (MSP) for identities
 - Administrator(s)
 - Users
 - Peers
 - Orderers (optional)
- A network can include many organisations representing a consortium
- Each organisation has an ID

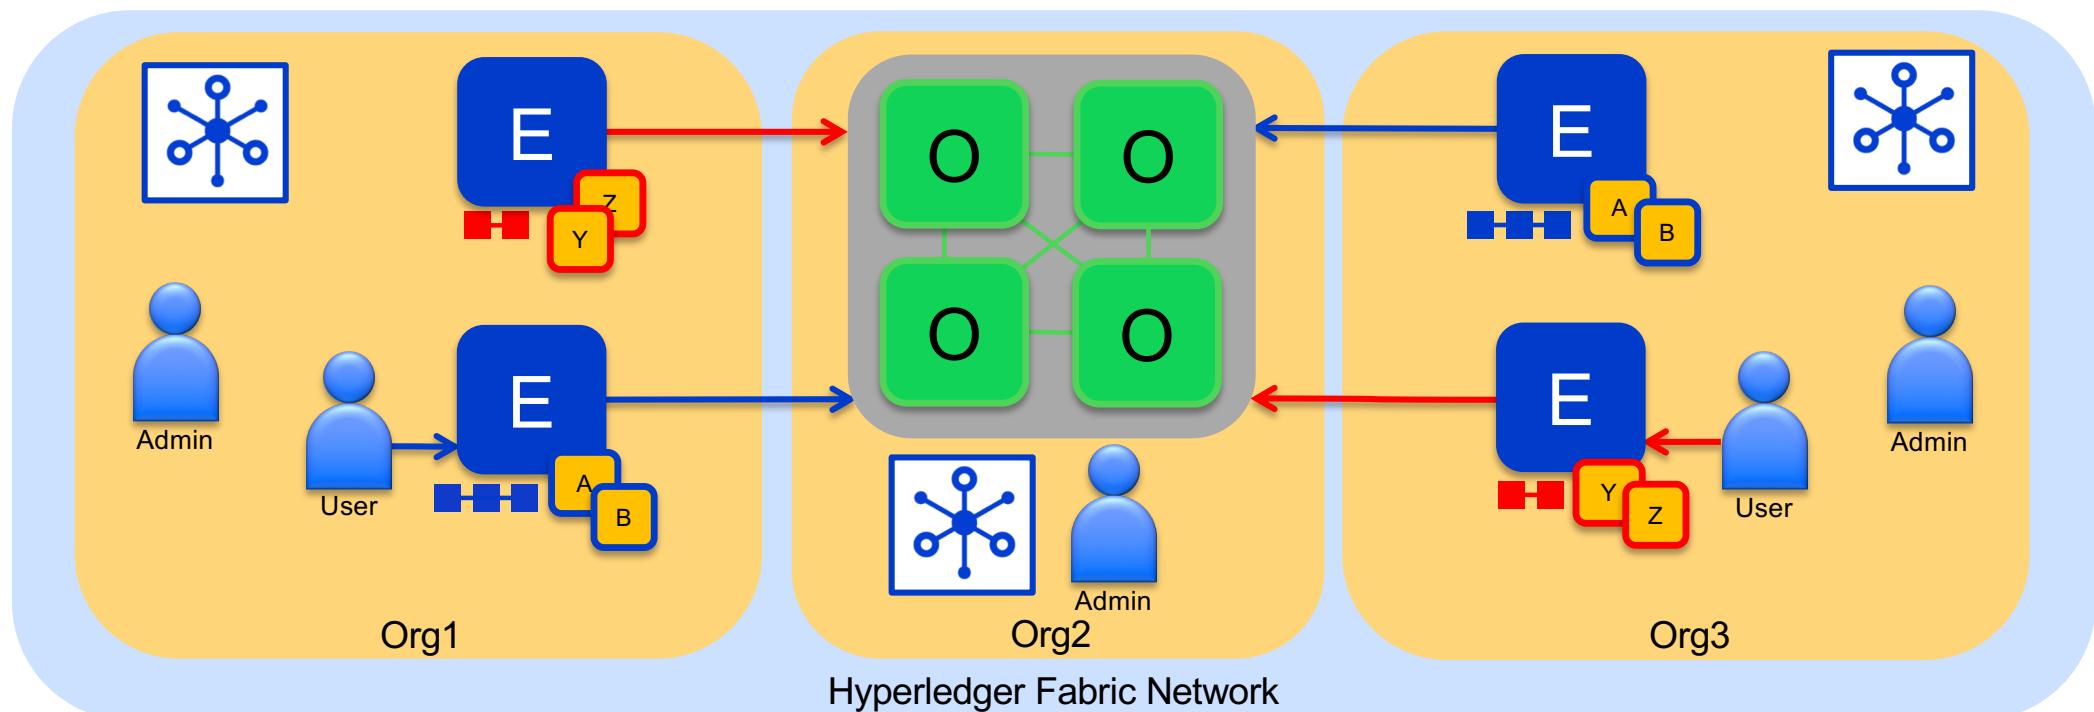


Consortium Network

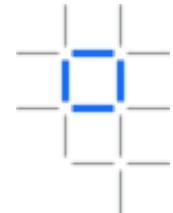


An example consortium network of 3 organisations

- Orgs 1 and 3 run peers
- Org 2 provides the ordering service only

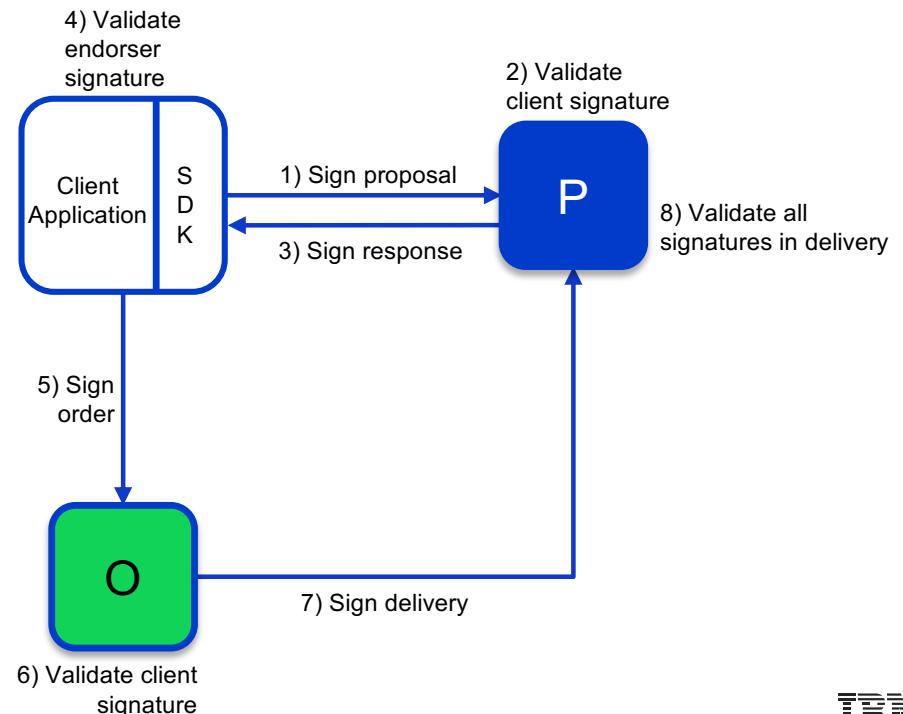


Transaction Signing

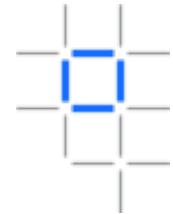


All transactions within a Hyperledger Fabric network are signed by permissioned actors, and those signatures validated

- Actors sign transactions with their enrolment private key
 - Stored in their local MSP
- Components validate transactions and certificates
 - Root CA certificates and CRLs stored in local MSP
 - Root CA certificates and CRLs stored in Org MSP in channel

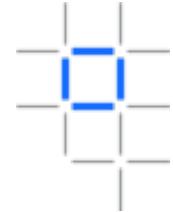


Summary and Next Steps



- Apply shared ledgers and smart contracts to your Business Network
- Think about your participants, assets and business processes
- Spend time thinking about realistic business use cases
- Get some hands-on experience with the technology
- Start with a First Project
- IBM can help with your journey

Further Hyperledger Fabric Information



- Project Home: <https://www.hyperledger.org/projects/fabric>
- GitHub Repo: <https://github.com/hyperledger/fabric>
- Latest Docs: <https://hyperledger-fabric.readthedocs.io/en/latest/>
- Community Chat: <https://chat.hyperledger.org/channel/fabric>
- Project Wiki: <https://wiki.hyperledger.org/projects/fabric>
- Design Docs: <https://wiki.hyperledger.org/community/fabric-design-docs>

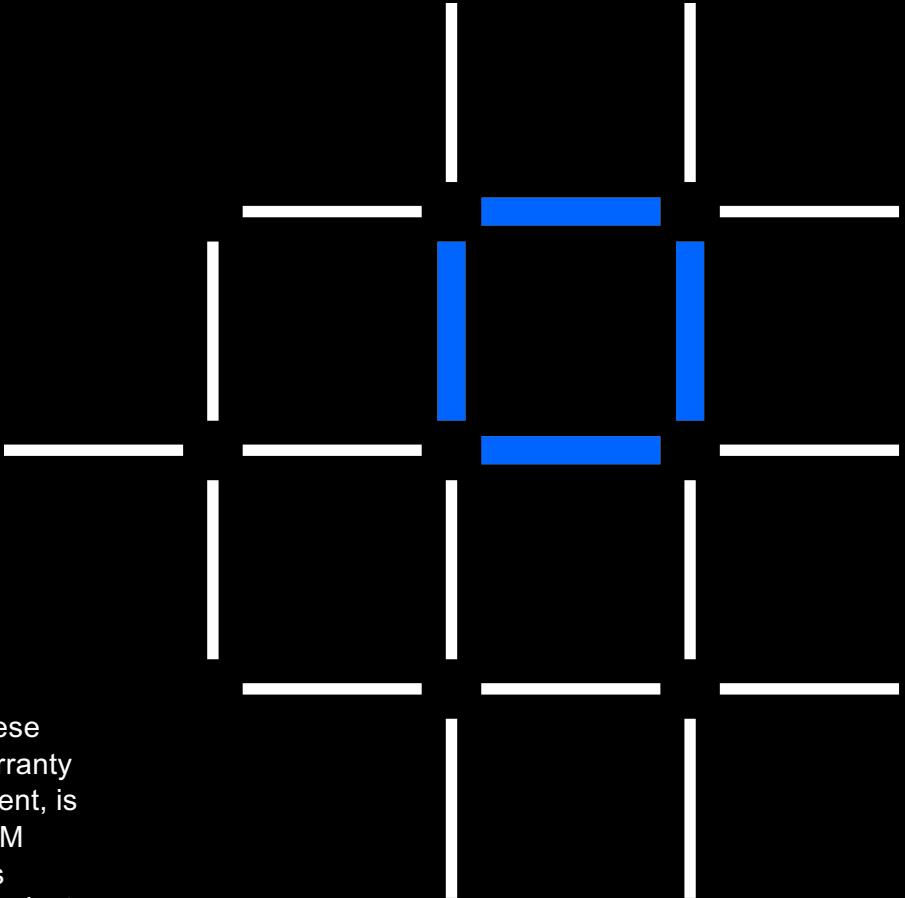
Thank you

IBM Blockchain

www.ibm.com/blockchain

developer.ibm.com/blockchain

www.hyperledger.org

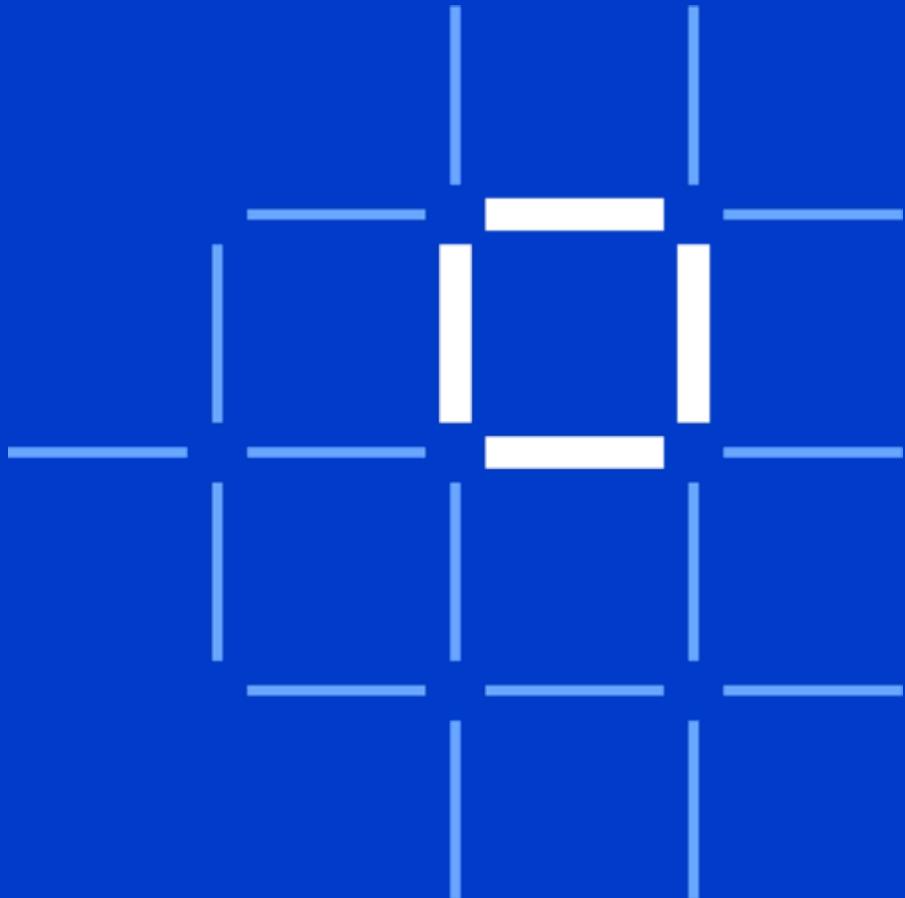


© Copyright IBM Corporation 2019. All rights reserved. The information contained in these materials is provided for informational purposes only, and is provided AS IS without warranty of any kind, express or implied. Any statement of direction represents IBM's current intent, is subject to change or withdrawal, and represents only goals and objectives. IBM, the IBM logo, and other IBM products and services are trademarks of the International Business Machines Corporation, in the United States, other countries or both. Other company, product, or service names may be trademarks or service marks of others.

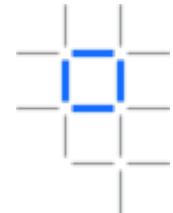
IBM

...

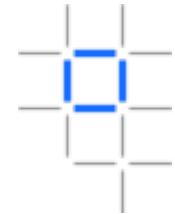
Appendix A: **Hyperledger Fabric Commands**



Fabric Commands



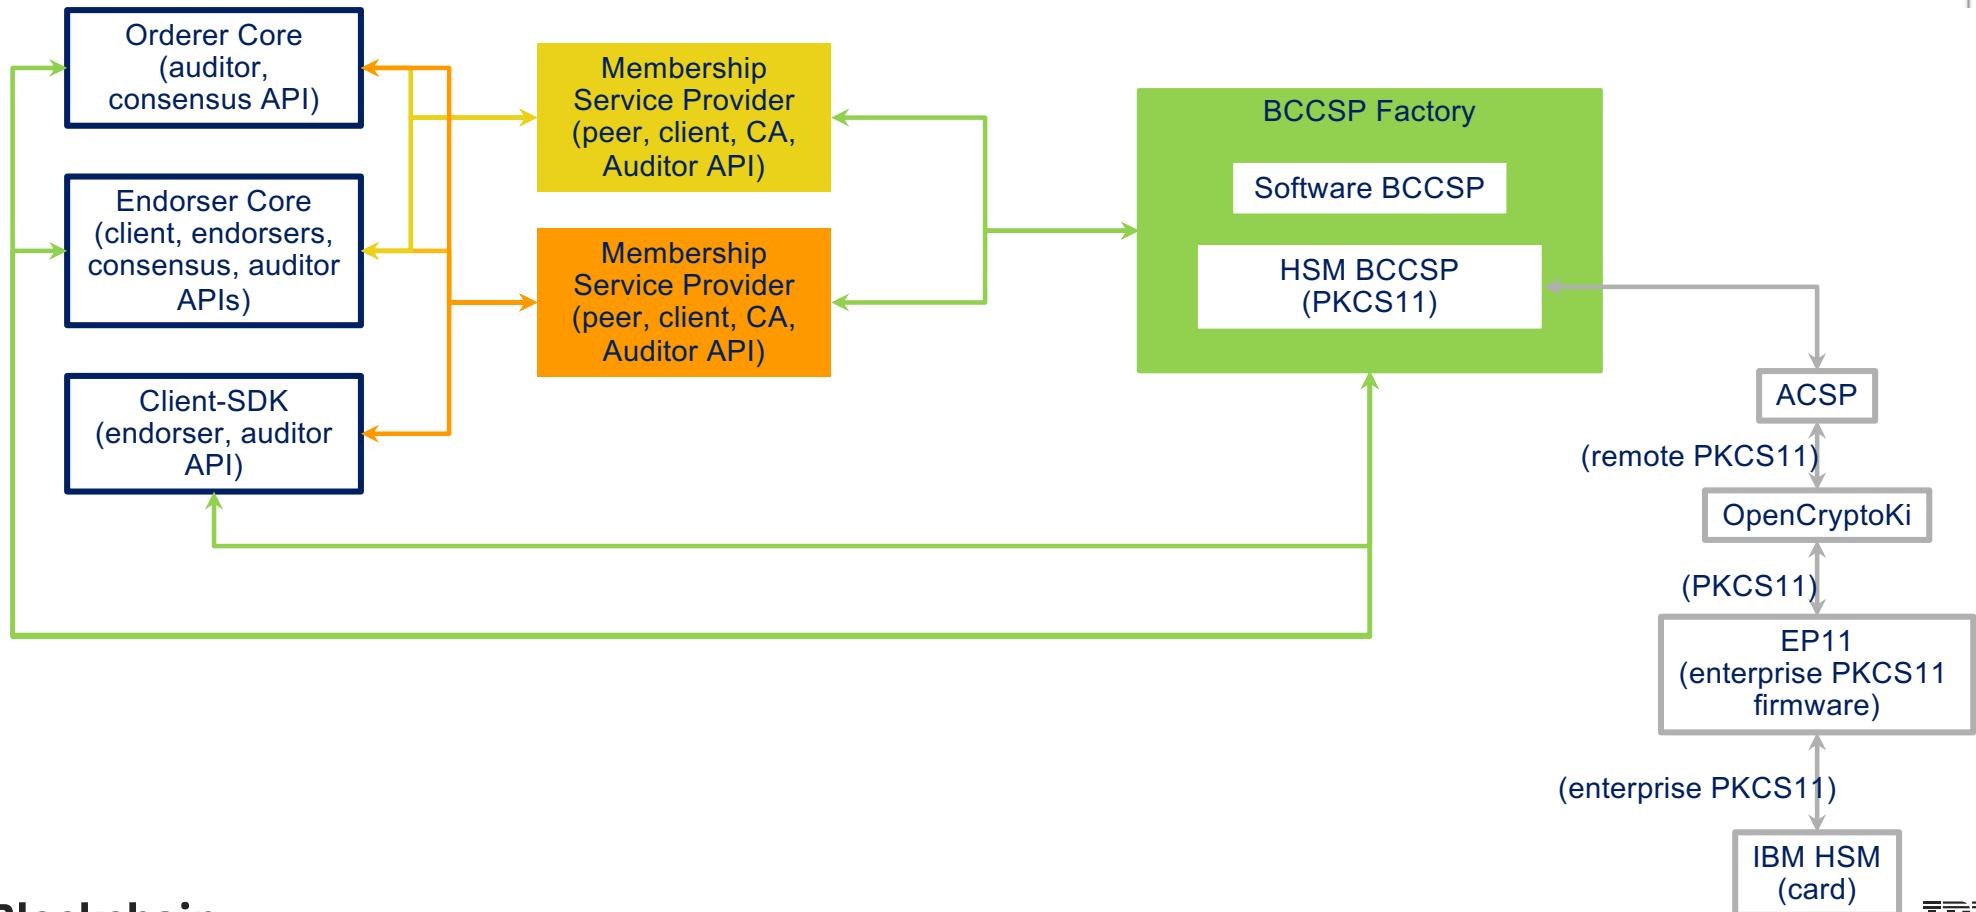
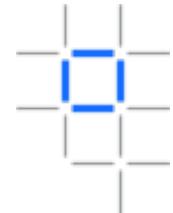
- Fabric has the following commands:
 - **peer** ... (For operating and configuring a peer)
 - **peer chaincode** ... (Manages chaincode on the peer)
 - **peer channel** ... (Manages channels on the peer)
 - **peer node** ... (Manages the peer)
 - **peer version** (Returns the peer version)
 - **cryptogen** ... (Utility for generating crypto material)
 - **configtxgen** ... (Creates configuration data such as the genesis block)
 - **configtxlator** ... (Utility for generating channel configurations)
 - **fabric-ca-client** ... (Manages identities)
 - **fabric-ca-server** ... (Manages the Fabric-CA server)



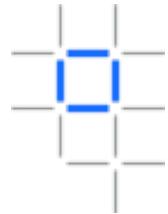
Configuration Detail

Path	MSP ID	Attributes	Attributes
config -> channel_group -> groups -> application -> groups	Org1MSP	mod_policy	Admins
		policies -> Admins	mod_policy
			Admins
			policy -> value -> identities
			Org1MSP, Admin
			Policy -> value -> rule
			n_out_of
		Policies -> Readers	Mod_policy
			Admins
			Policy -> value -> identities
			Org1MSP
			Policy -> value -> rule
			N_out_of
		Policies -> Writers	Mod_policy
			Admins
			Policy -> value -> identities
			Org1MSP
			Policy -> value -> rule
			N_out_of

MSP and BCCSP (Modularity and Decentralisation)



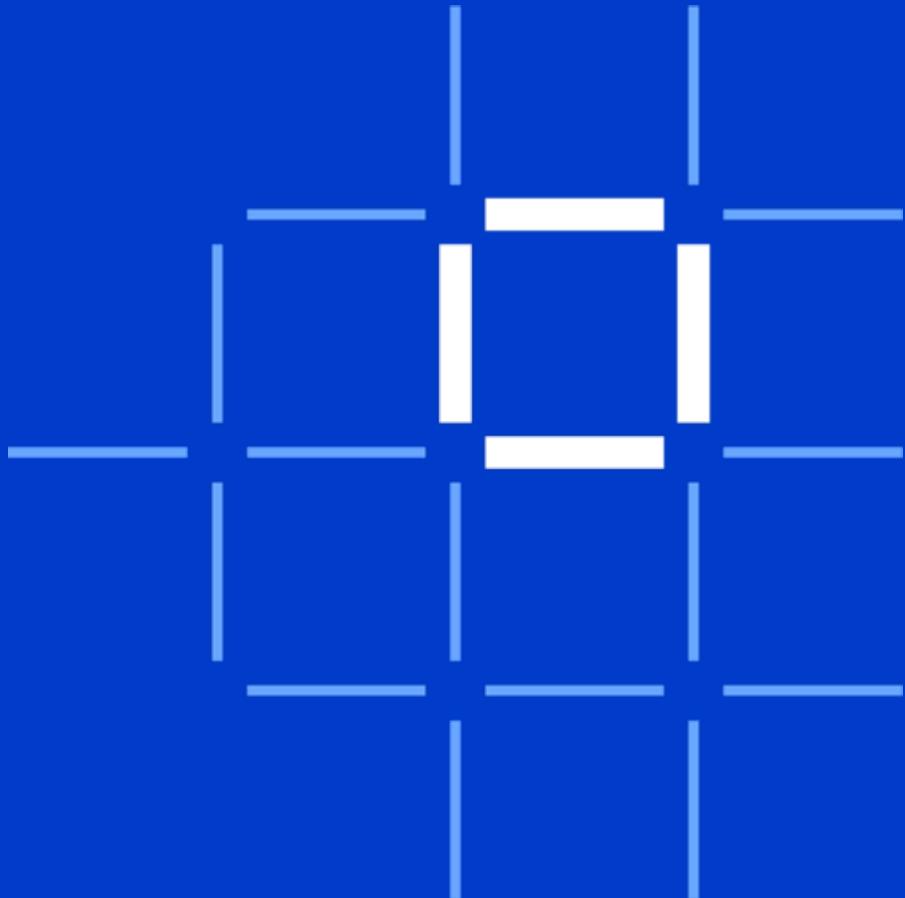
Blockchain Crypto Service Provider (BCCSP)

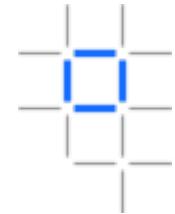


- Pluggable implementation of cryptographic standards and algorithms.
- **Pluggability**
 - alternate implementations of crypto interface can be used within the Hyperledger Fabric code, without modifying the core
- **Support for Multiple CSPs**
 - Easy addition of more types of CSPs, e.g., of different HSM types
 - Enable the use of different CSP on different system components transparently
- **International Standards Support**
 - E.g., via a new/separate CSP
 - Interoperability among standards is not necessarily guaranteed

...

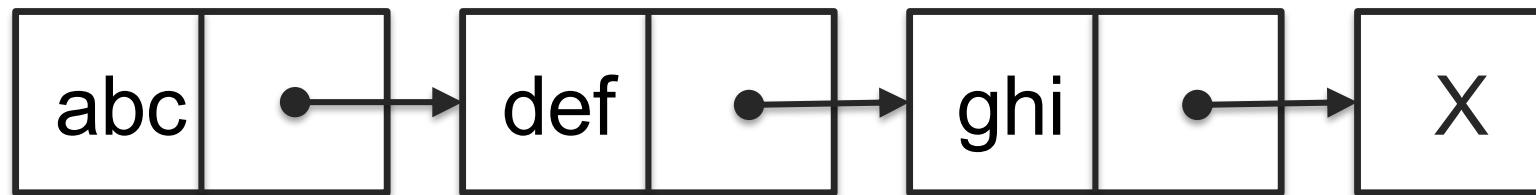
Appendix B: **Blockchain Data Structures Overview**

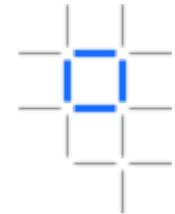




The Linked List

- Linear collection of data elements
- Each element is linked to the next
- Concept dates from 1955

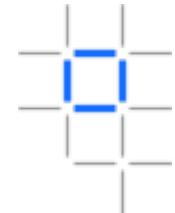




One-Way Hash Functions

- Any function that can be applied to a set of data that is guaranteed to produce the same output for the same input
- One-way means that you can't derive the input from the output
- Often, outputs are *unlikely* to repeat for different inputs
- Forms the basis of much cryptography

$$\begin{aligned} h(abc) &= 7859 \\ h(def) &= 8693 \\ h'(7859) &= ? \\ h(abc) &= 7859 \end{aligned}$$



The Hash Chain

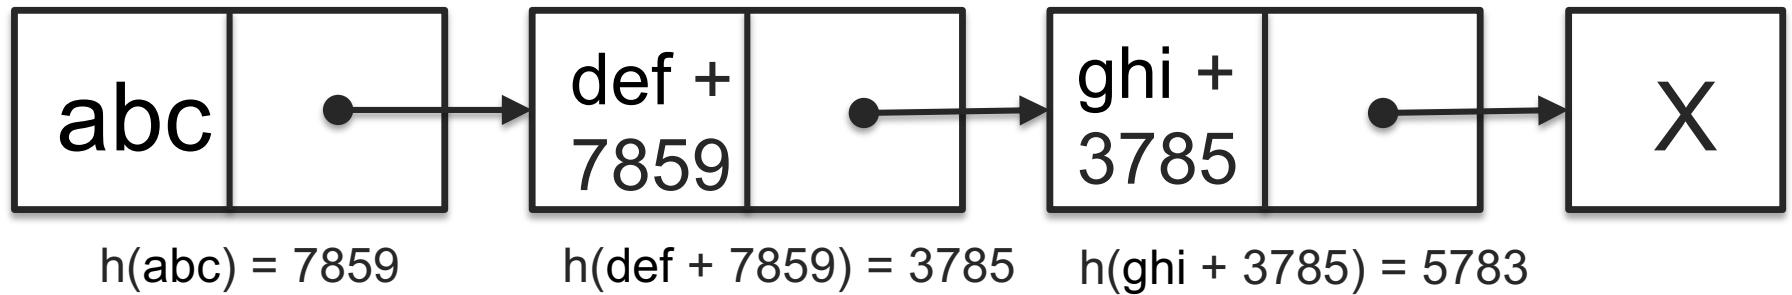
Hash chain: A successive application of a hash function

$$h(h(h(abc))) = 1859$$

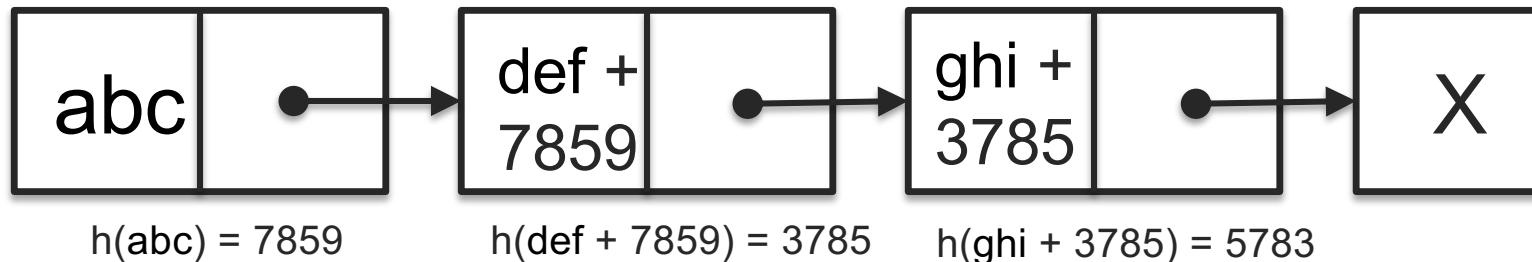
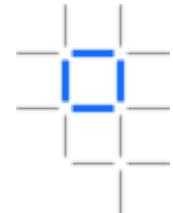
Can combine new data with each successive hash:

$$h(ghi + h(def + h(abc))) = 5783$$

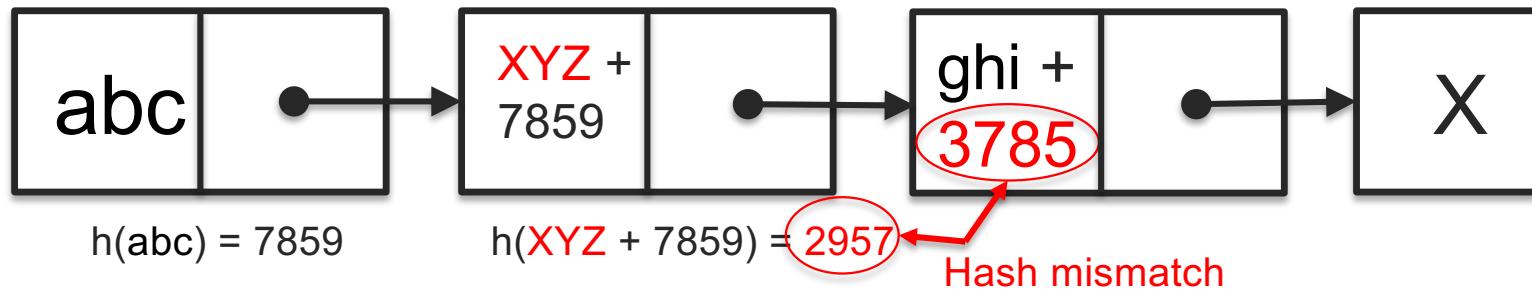
Using this concept you can produce a tamper resistant linked list:

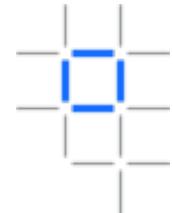


How is This Tamper Resistant?



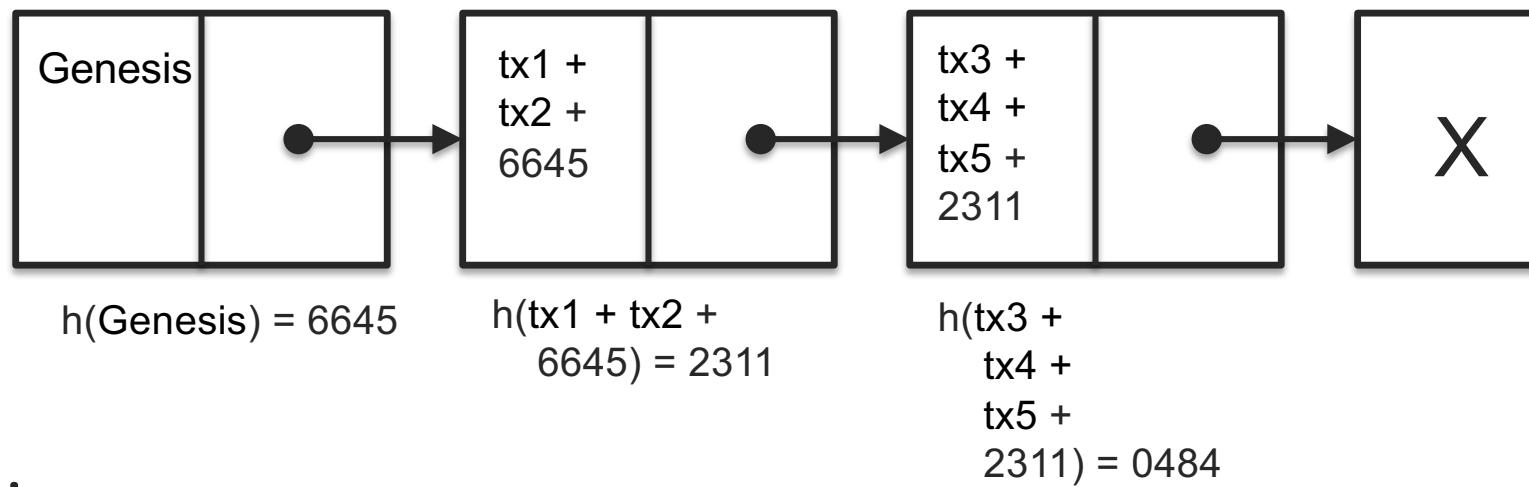
- Any modification to a data element means that the hashes will not match up
 - You would need to recreate the downstream chain



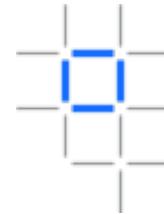


Applied to Blockchain

- A blockchain is a hash chain (*with optimizations* that we'll cover shortly)
- Each element (block) in the linked list is a set of zero or more transactions
 - Transactions are an implementation-dependent data object
- First block known as a genesis block
 - May contain some identifying string or other configuration metadata

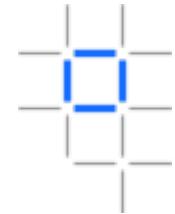


Some Problems with This Approach

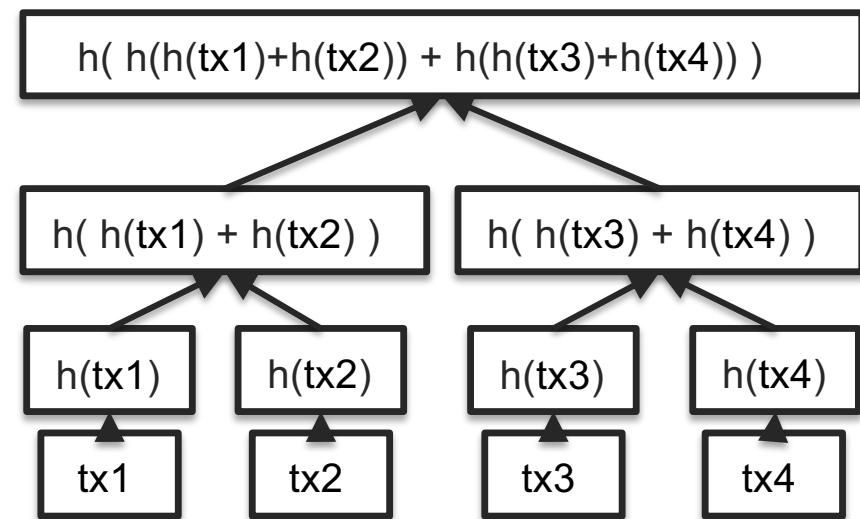


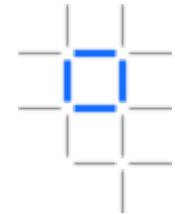
- In the event of tampering, it can be difficult to identify which transaction was modified (particularly when there are many transactions in a block)
 - It is not feasible to have one transaction per block
- It requires all transaction data in order to retain integrity of chain
- Searching transactions is linear (time consuming)

Merkle Tree



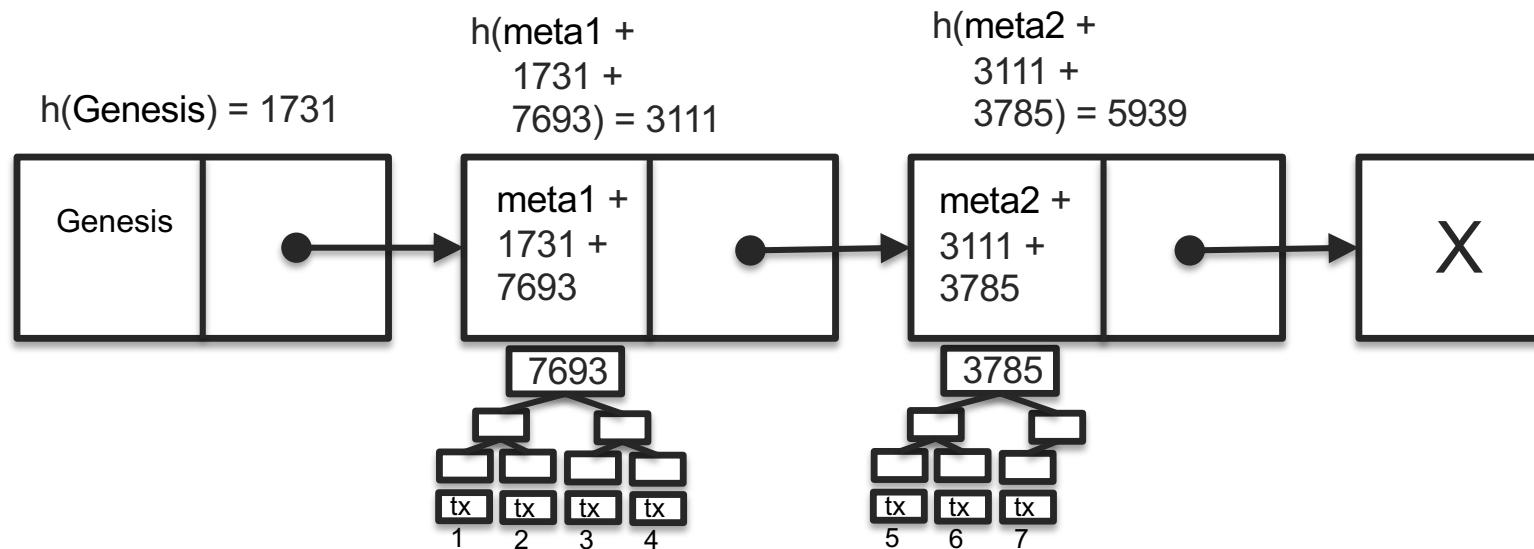
- It is possible to optimise the chain data structure if we arrange it as a tree
 - Makes it easier to identify tampering without sacrificing stability
 - Makes it quicker to traverse
- However, this makes it impossible to add new transactions without re-hashing root nodes



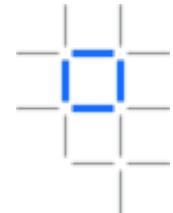


Combining Tree and Chain

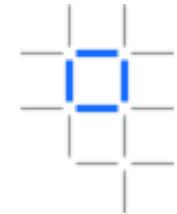
- Each element in the chain contains:
 - A pointer (“Merkle root”) to the tree of transactions
 - Other metadata (e.g. timestamp)
 - A hash of the previous block’s data (i.e. Merkle root, metadata and hash)



Benefits of This Approach



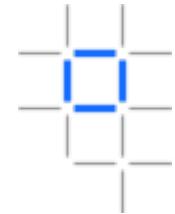
- It allows tampered transactions to be identified easily
- More [efficient search](#) within a block
 - $O(\log N)$ rather than $O(N)$
- Allows transaction detail to be stubbed
 - Bitcoin has a [Simplified Payment Verifier](#) (SPV) concept: a type of user that doesn't have the entire tree available, just the Merkel roots
 - Note it is also possible to checkpoint and archive old blocks, creating a new Genesis block mid-way through the chain



Common Transactions

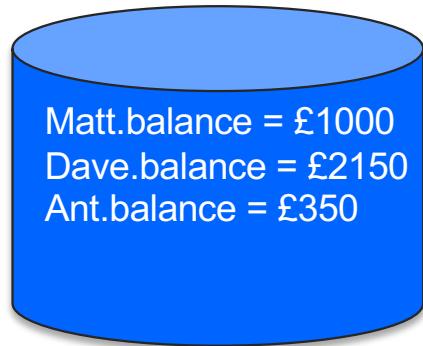
- What's Dave's balance?
- Does Matt have funds to clear a £1000 transaction? (Assuming no overdraft)

#	Transaction	Initiator	Receiver	Amount
1	Create a/c	Cash	Matt	£1000
2	Create a/c	Cash	Dave	£2000
3	Transfer	Matt	Dave	£100
4	Create a/c	Cash	Ant	£500
5	Transfer	Ant	Matt	£50
6	Transfer	Ant	Dave	£200
7	Transfer	Dave	Matt	£100
8	Transfer	Dave	Ant	£50
9	Transfer	Matt	Ant	£50

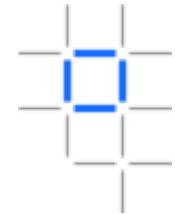


World State

- It is clearly not feasible to reparse the entire transaction log to complete a new transaction
- Blockchains often include an associated database (world state) – e.g. Hyperledger Fabric
- Transactions become a set of creates, reads, updates and deletes of records in this data store



#	Transaction	Initiator	Receiver	Amount
1	Create a/c	Cash	Matt	£1000
2	Create a/c	Cash	Dave	£2000
3	Transfer	Matt	Dave	£100
4	Create a/c	Cash	Ant	£500
5	Transfer	Ant	Matt	£50
6	Transfer	Ant	Dave	£200
7	Transfer	Dave	Matt	£100
8	Transfer	Dave	Ant	£50
9	Transfer	Matt	Ant	£50



Unspent Transaction Outputs

- Some blockchains (e.g. Bitcoin) don't maintain balances
 - Transactions are linked to earlier transactions using an ID (TXID)
 - Outputs always equal inputs
 - Unspent funds are marked as an “Unspent Transaction Output” (UTXO)
 - Only UTXOs can be used as inputs (to prevent double spending)
 - Your “balance” is the aggregation of all of your UTXOs
- In Bitcoin, if your application doesn't specify the UTXO output then the miner gets the excess!

